

Debarghya Dey Enrollment No.- 510519087 Assignment - 6

Project Setup

Mount Google Drive

```
In [7]: from google.colab import drive  
drive.mount('/content/drive')  
  
Mounted at /content/drive
```

Importing all necessary libraries

```
In [8]: # system info  
import platform  
import random  
import os  
import sys  
from collections import OrderedDict  
  
# data visualisation and manipulation  
import numpy as np  
import pandas as pd  
import matplotlib  
import matplotlib.pyplot as plt  
from matplotlib import style  
import seaborn as sns  
import sklearn  
from sklearn.metrics import classification_report, precision_score, recall_score, f1_score, precision_recall_fscore_support  
  
#model selection  
from sklearn.model_selection import train_test_split  
from sklearn.model_selection import KFold  
from sklearn.metrics import accuracy_score,precision_score,recall_score,confusion_matrix,roc_curve,roc_auc_score  
from sklearn.model_selection import GridSearchCV  
from sklearn.preprocessing import LabelEncoder  
  
#preprocess.
```

```
from keras.preprocessing.image import ImageDataGenerator

#dl libraries
from keras import backend as K
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam,SGD,Adagrad,Adadelta,RMSprop
from keras.utils import to_categorical

# specifically for cnn
from keras.layers import Dropout, Flatten, Activation
from keras.layers import Conv2D, MaxPooling2D, BatchNormalization, AveragePooling2D
import tensorflow as tf
import random as rn
import keras

# specifically for manipulating zipped images and getting numpy arrays of pixel values of images.
import cv2
import numpy as np
from tqdm import tqdm
import os
from random import shuffle
from zipfile import ZipFile
from PIL import Image
```

Data-Loading

Create training data

```
In [9]: X = []
Z = []
IMG_SIZE = 80
BASE_PATH = "/content/drive/My Drive/ML_DRIVE/Assign_6/flowers_dataset/"
FLOWER_DAISY_DIR = BASE_PATH + "daisy"
FLOWER_SUNFLOWER_DIR = BASE_PATH + "sunflower"
FLOWER_TULIP_DIR = BASE_PATH + "tulip"
FLOWER_DANDI_DIR = BASE_PATH + "dandelion"
FLOWER_ROSE_DIR = BASE_PATH + "rose"
```

```
In [10]: def assign_label(img, flower_type):
    return flower_type

def make_train_data(flower_type, DIR, mode = 'color'):
    for img in tqdm(os.listdir(DIR)):
        label = assign_label(img, flower_type)
        path = os.path.join(DIR, img)
        if (mode == 'color'):
            img = cv2.imread(path, cv2.IMREAD_COLOR)
        else:
            img = cv2.imread(path, cv2.IMREAD_GRAYSCALE)

        img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))

        X.append(img)
        Z.append(str(label))
    return
```

```
In [11]: make_train_data('Daisy', FLOWER_DAISY_DIR, 'gray')
```

```
100%|██████████| 764/764 [00:13<00:00, 54.68it/s]
```

```
In [12]: make_train_data('Sunflower', FLOWER_SUNFLOWER_DIR, 'gray')
```

```
100%|██████████| 733/733 [00:11<00:00, 63.90it/s]
```

```
In [13]: make_train_data('Tulip', FLOWER_TULIP_DIR, 'gray')
```

```
100%|██████████| 984/984 [00:22<00:00, 44.56it/s]
```

```
In [14]: make_train_data('Dandelion', FLOWER_DANDI_DIR, 'gray')
```

```
100%|██████████| 1052/1052 [00:18<00:00, 57.20it/s]
```

```
In [15]: make_train_data('Rose', FLOWER_ROSE_DIR, 'gray')
```

```
100%|██████████| 784/784 [00:11<00:00, 70.37it/s]
```

Define helper functions for plotting accuracy, loss, confusion matrix for all models

```
In [16]: def draw_graph(history):
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Model Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['Training Accuracy', 'Validation Accuracy'], loc='upper left')
    plt.show()
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Loss vs Epoch')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['Training Loss', 'Validation Loss'], loc='upper right')
    plt.show()

    print("Maximum training accuracy: ", max(history.history['accuracy']))
    print("Maximum validation accuracy: ", max(history.history['val_accuracy']))

def draw_matrix(model, train, test, train_labels, test_labels):
    predictions_test = []
    for item in test:
        predictions_test.append(np.argmax(model.predict(item.reshape(-1,80,80,1)))))

    cm_test = sklearn.metrics.confusion_matrix(np.argmax(test_labels, axis = 1), predictions_test)
    sns.heatmap(cm_test, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues_r');
    plt.ylabel('Actual label');
    plt.xlabel('Predicted label');
    all_sample_title = 'Confusion Matrix (Test samples)'
    plt.title(all_sample_title)
    plt.show()
    print(classification_report(np.argmax(test_labels, axis = 1), predictions_test))

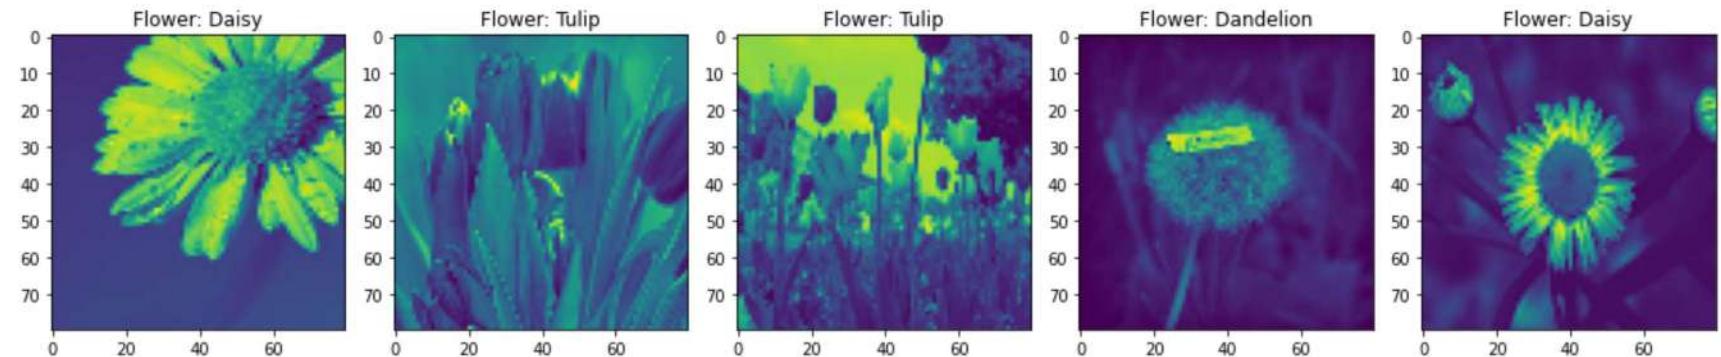
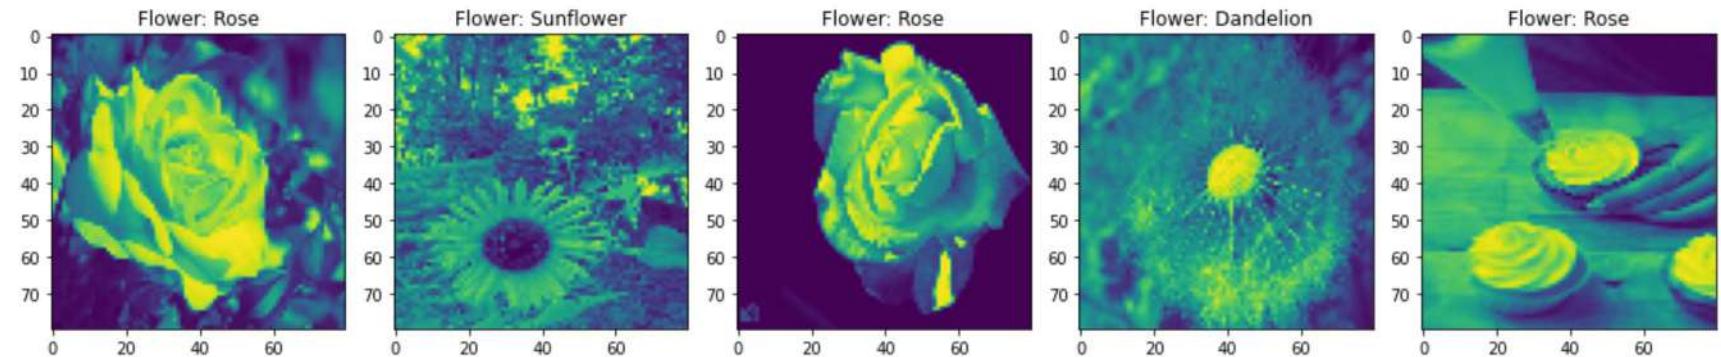
def draw(History, model, train, test, train_labels, test_labels):
    draw_graph(History)
    draw_matrix(model, train, test, train_labels, test_labels)
```

Visualizing some of the images we loaded

```
In [17]: fig,ax = plt.subplots(2,5)
fig.set_size_inches(15,15)
```

```
for i in range(2):
    for j in range (5):
        l = rn.randint(0, len(Z))
        ax[i,j].imshow(X[l])
        ax[i,j].set_title('Flower: '+ Z[l])

plt.tight_layout()
```



Pre-processing

One-Hot Encoding

```
In [18]: le = LabelEncoder()
Y = le.fit_transform(Z)
Y = to_categorical(Y,5)
X = np.array(X)
X = X / 255
```

Train : Test split

```
In [19]: train, test, train_labels, test_labels = train_test_split(X, Y, test_size=0.1, random_state=100)
train = train.reshape((-1,80,80,1))
test = test.reshape((-1,80,80,1))
```

Trying out the initial 4 models

Model 1

```
In [24]: model1 = Sequential()

model1.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model1.add(MaxPooling2D(pool_size=(2, 2)))
model1.add(Dropout(0.1))

model1.add(Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',activation ='relu'))
model1.add(MaxPooling2D(pool_size=(2, 2)))
model1.add(Dropout(0.1))

model1.add(Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',activation ='relu'))
model1.add(MaxPooling2D(pool_size=(2, 2)))
model1.add(Dropout(0.1))

model1.add(Flatten())
model1.add(Dense(5, activation = 'softmax'))
```

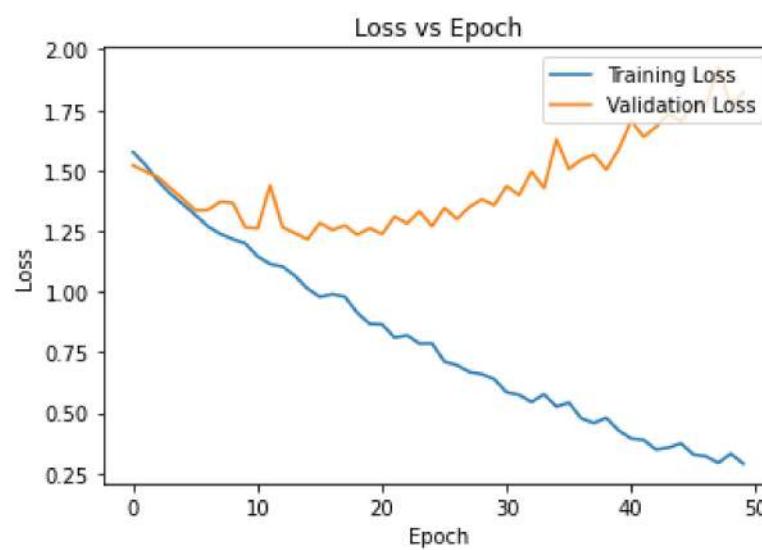
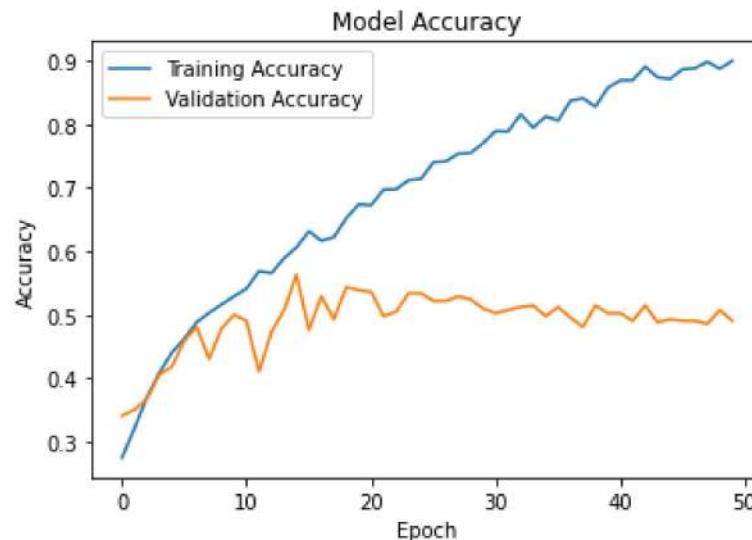
```
model1.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])  
model1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d (MaxPooling2D)	(None, 40, 40, 16)	0
dropout (Dropout)	(None, 40, 40, 16)	0
conv2d_1 (Conv2D)	(None, 40, 40, 32)	4640
max_pooling2d_1 (MaxPooling 2D)	(None, 20, 20, 32)	0
dropout_1 (Dropout)	(None, 20, 20, 32)	0
conv2d_2 (Conv2D)	(None, 20, 20, 64)	18496
max_pooling2d_2 (MaxPooling 2D)	(None, 10, 10, 64)	0
dropout_2 (Dropout)	(None, 10, 10, 64)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 5)	32005
<hr/>		
Total params: 55,301		
Trainable params: 55,301		
Non-trainable params: 0		

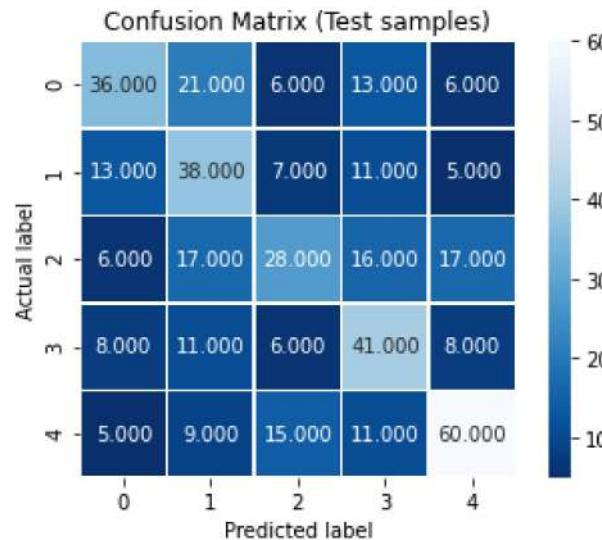
```
In [ ]: History1 = model1.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

```
In [ ]: draw(History1, model1, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.8999999761581421

Maximum validation accuracy: 0.5628019571304321



	precision	recall	f1-score	support
0	0.53	0.44	0.48	82
1	0.40	0.51	0.45	74
2	0.45	0.33	0.38	84
3	0.45	0.55	0.49	74
4	0.62	0.60	0.61	100
accuracy			0.49	414
macro avg	0.49	0.49	0.48	414
weighted avg	0.50	0.49	0.49	414

Model 2

```
In [ ]: model2 = Sequential()

model2.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model2.add(MaxPooling2D(pool_size=(2, 2)))
model2.add(Dropout(0.1))

model2.add(Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',activation ='relu'))
model2.add(MaxPooling2D(pool_size=(2, 2)))
model2.add(Dropout(0.1))

model2.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
```

```
model2.add(MaxPooling2D(pool_size=(2, 2)))
model2.add(Dropout(0.1))

model2.add(Flatten())
model2.add(Dense(5, activation = 'softmax'))

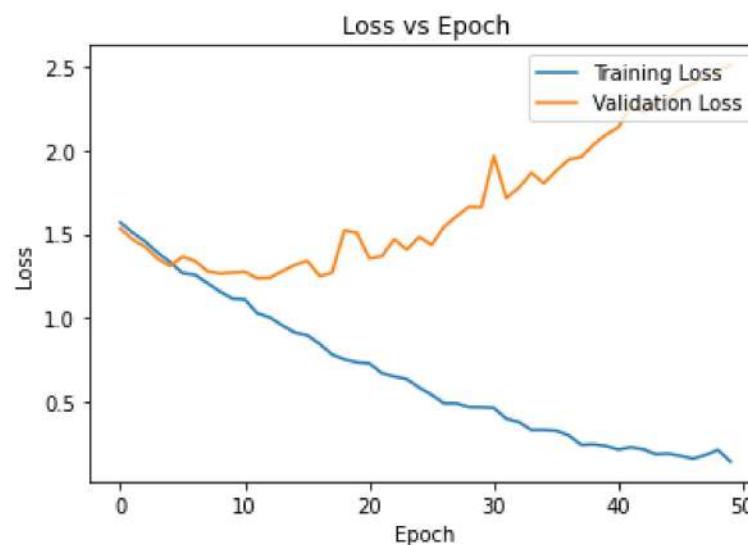
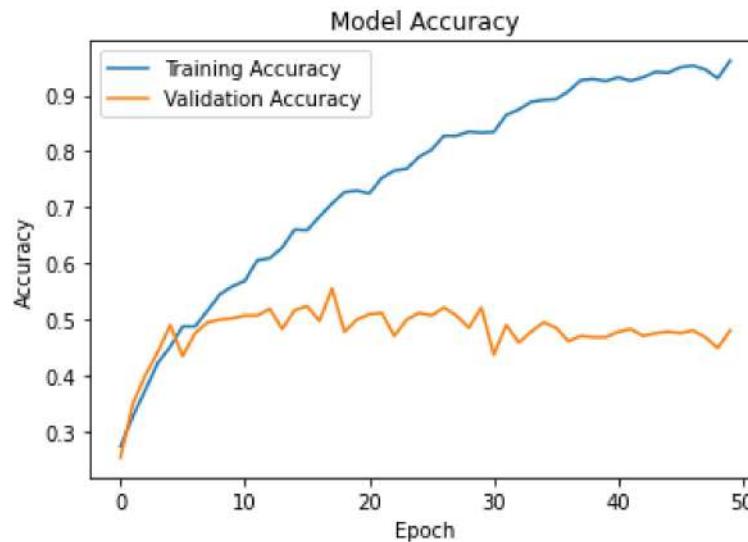
model2.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model2.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_3 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
max_pooling2d_3 (MaxPooling2	(None, 40, 40, 16)	0
<hr/>		
dropout_3 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_4 (Conv2D)	(None, 40, 40, 32)	4640
<hr/>		
max_pooling2d_4 (MaxPooling2	(None, 20, 20, 32)	0
<hr/>		
dropout_4 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_5 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_5 (MaxPooling2	(None, 10, 10, 64)	0
<hr/>		
dropout_5 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_1 (Flatten)	(None, 6400)	0
<hr/>		
dense_1 (Dense)	(None, 5)	32005
<hr/>		
Total params: 88,069		
Trainable params: 88,069		
Non-trainable params: 0		

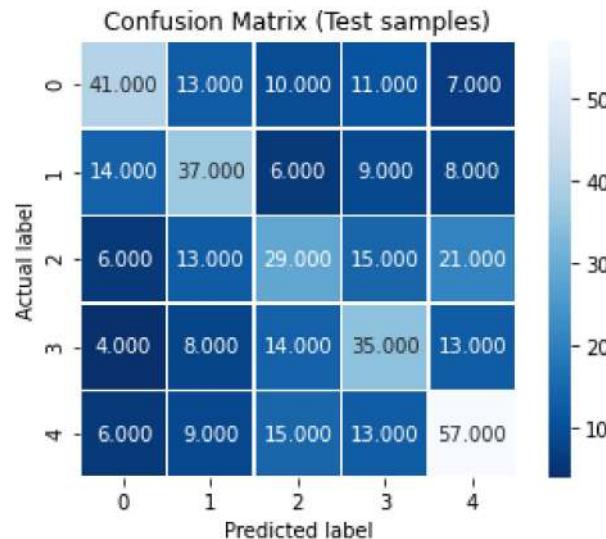
In []: History2 = model2.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History2, model2, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9620967507362366

Maximum validation accuracy: 0.5555555820465088



	precision	recall	f1-score	support
0	0.58	0.50	0.54	82
1	0.46	0.50	0.48	74
2	0.39	0.35	0.37	84
3	0.42	0.47	0.45	74
4	0.54	0.57	0.55	100
accuracy			0.48	414
macro avg	0.48	0.48	0.48	414
weighted avg	0.48	0.48	0.48	414

Model 3

```
In [ ]: model3 = Sequential()

model3.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))

model3.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))

model3.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
```

```
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))

model3.add(Flatten())
model3.add(Dense(5, activation = 'softmax'))

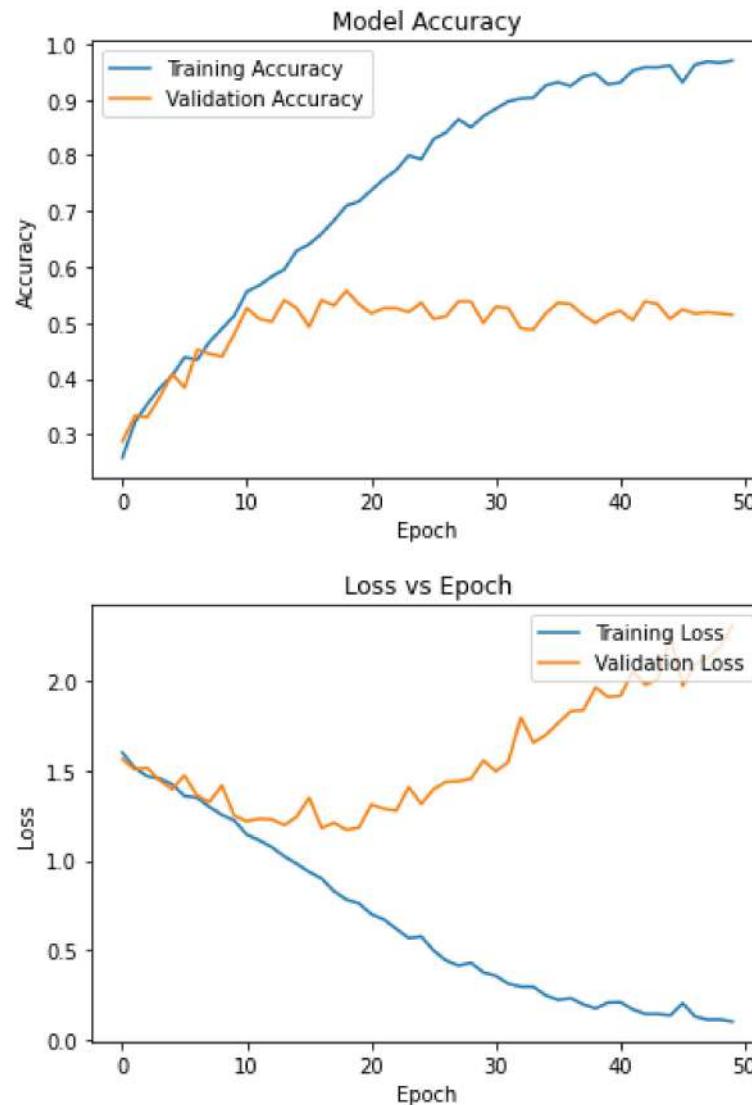
model3.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model3.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_6 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
max_pooling2d_6 (MaxPooling2	(None, 40, 40, 16)	0
<hr/>		
dropout_6 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_7 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
max_pooling2d_7 (MaxPooling2	(None, 20, 20, 32)	0
<hr/>		
dropout_7 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_8 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_8 (MaxPooling2	(None, 10, 10, 64)	0
<hr/>		
dropout_8 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_2 (Flatten)	(None, 6400)	0
<hr/>		
dense_2 (Dense)	(None, 5)	32005
<hr/>		
Total params: 96,261		
Trainable params: 96,261		
Non-trainable params: 0		

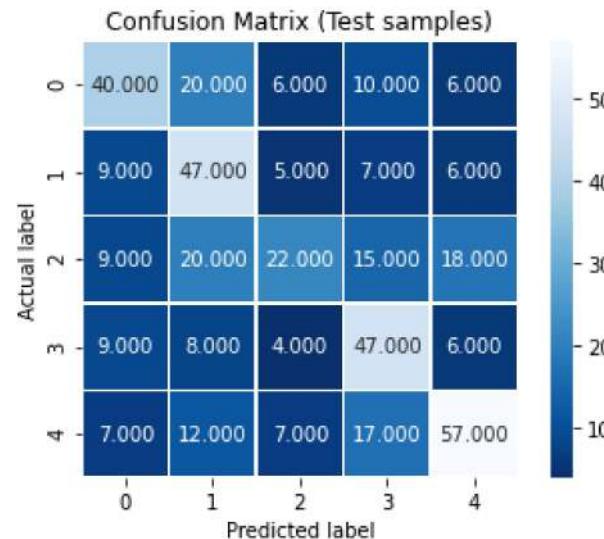
In []: History3 = model3.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History3, model3, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9701613187789917

Maximum validation accuracy: 0.5579710006713867



	precision	recall	f1-score	support
0	0.54	0.49	0.51	82
1	0.44	0.64	0.52	74
2	0.50	0.26	0.34	84
3	0.49	0.64	0.55	74
4	0.61	0.57	0.59	100
accuracy			0.51	414
macro avg	0.52	0.52	0.50	414
weighted avg	0.52	0.51	0.51	414

Model 4

```
In [ ]: model4 = Sequential()

model4.add(Conv2D(filters = 16, kernel_size = (5,5),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model4.add(MaxPooling2D(pool_size=(2, 2)))
model4.add(Dropout(0.1))

model4.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model4.add(MaxPooling2D(pool_size=(2, 2)))
model4.add(Dropout(0.1))

model4.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
```

```

model4.add(MaxPooling2D(pool_size=(2, 2)))
model4.add(Dropout(0.1))

model4.add(Flatten())
model4.add(Dense(5, activation = 'softmax'))

model4.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model4.summary()

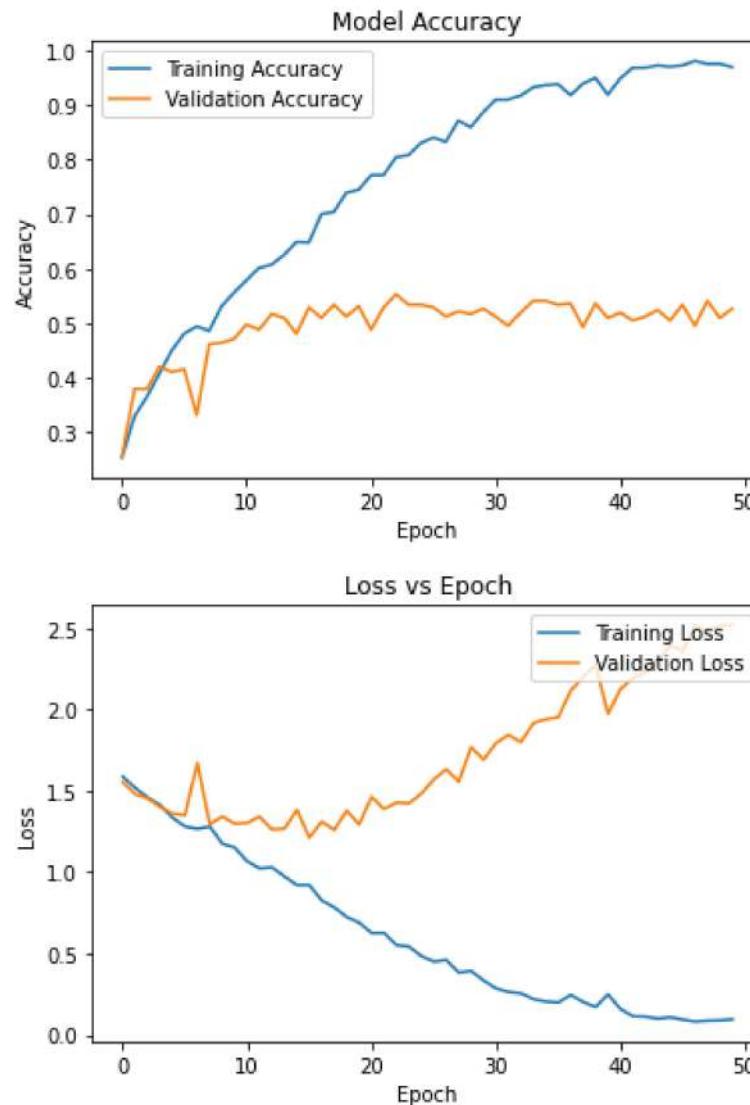
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_9 (Conv2D)	(None, 80, 80, 16)	416
<hr/>		
max_pooling2d_9 (MaxPooling2D)	(None, 40, 40, 16)	0
<hr/>		
dropout_9 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_10 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
max_pooling2d_10 (MaxPooling2D)	(None, 20, 20, 32)	0
<hr/>		
dropout_10 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_11 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_11 (MaxPooling2D)	(None, 10, 10, 64)	0
<hr/>		
dropout_11 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_3 (Flatten)	(None, 6400)	0
<hr/>		
dense_3 (Dense)	(None, 5)	32005
<hr/>		
Total params: 96,517		
Trainable params: 96,517		
Non-trainable params: 0		

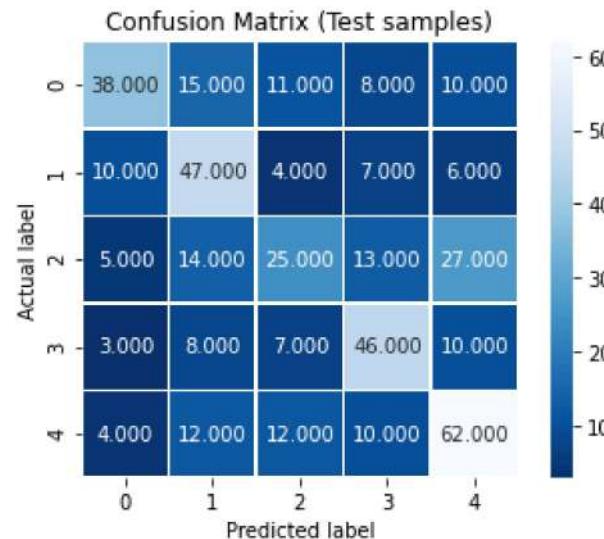
In []: History4 = model4.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History4, model4, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9814516305923462

Maximum validation accuracy: 0.5531401038169861



	precision	recall	f1-score	support
0	0.63	0.46	0.54	82
1	0.49	0.64	0.55	74
2	0.42	0.30	0.35	84
3	0.55	0.62	0.58	74
4	0.54	0.62	0.58	100
accuracy			0.53	414
macro avg	0.53	0.53	0.52	414
weighted avg	0.53	0.53	0.52	414

Model 3 was the best model with an accuracy of 57.84% on the validation dataset.

The model is as follows:

```
model3 = Sequential()

model3.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))
```

```
model3.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))

model3.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model3.add(MaxPooling2D(pool_size=(2, 2)))
model3.add(Dropout(0.1))

model3.add(Flatten())
model3.add(Dense(5, activation = 'softmax'))

model3.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=[ 'accuracy'])
model3.summary()
```

Now, we will be adding 2/3 fully-connected layers to this model.

Model 5

```
In [ ]: model5 = Sequential()

model5.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model5.add(MaxPooling2D(pool_size=(2, 2)))
model5.add(Dropout(0.1))

model5.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model5.add(MaxPooling2D(pool_size=(2, 2)))
model5.add(Dropout(0.1))

model5.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model5.add(MaxPooling2D(pool_size=(2, 2)))
model5.add(Dropout(0.1))

model5.add(Flatten())
model5.add(Dense(128, activation = 'relu'))
model5.add(Dense(5, activation = 'softmax'))
```

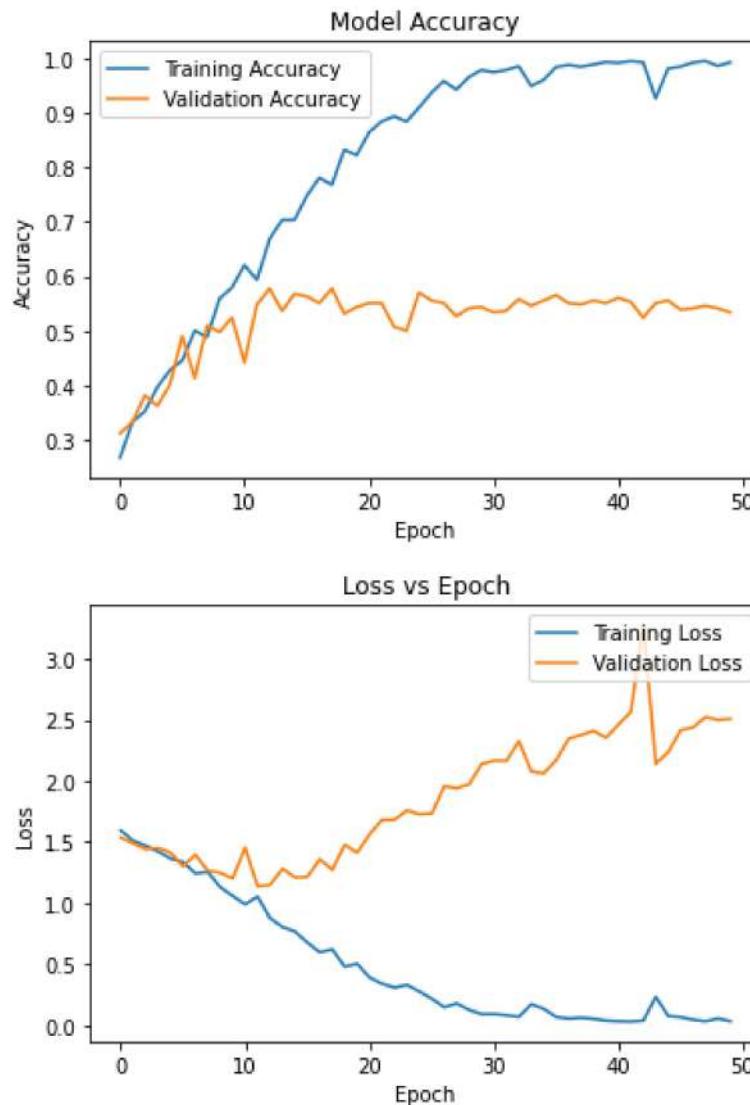
```
model5.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])  
model5.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_12 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
max_pooling2d_12 (MaxPooling)	(None, 40, 40, 16)	0
<hr/>		
dropout_12 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_13 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
max_pooling2d_13 (MaxPooling)	(None, 20, 20, 32)	0
<hr/>		
dropout_13 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_14 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_14 (MaxPooling)	(None, 10, 10, 64)	0
<hr/>		
dropout_14 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_4 (Flatten)	(None, 6400)	0
<hr/>		
dense_4 (Dense)	(None, 128)	819328
<hr/>		
dense_5 (Dense)	(None, 5)	645
<hr/>		
Total params: 884,229		
Trainable params: 884,229		
Non-trainable params: 0		

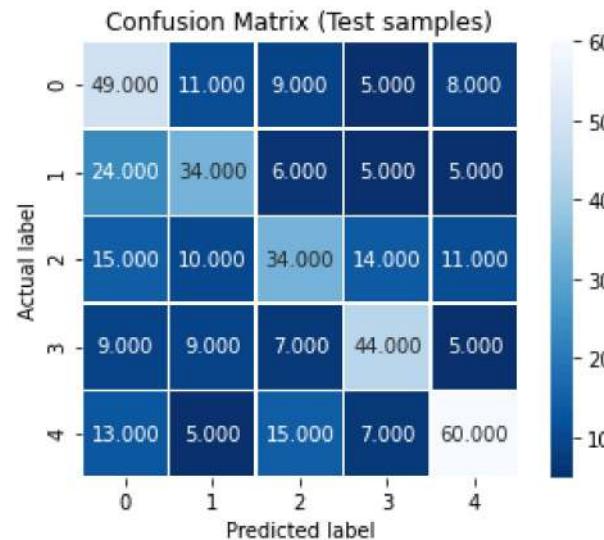
In []: History5 = model5.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History5, model5, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9946236610412598

Maximum validation accuracy: 0.5772947072982788



	precision	recall	f1-score	support
0	0.45	0.60	0.51	82
1	0.49	0.46	0.48	74
2	0.48	0.40	0.44	84
3	0.59	0.59	0.59	74
4	0.67	0.60	0.63	100
accuracy			0.53	414
macro avg	0.54	0.53	0.53	414
weighted avg	0.54	0.53	0.53	414

Model 6

```
In [ ]: model6 = Sequential()

model6.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))

model6.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))
model6.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model6.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model6.add(Dropout(0.1))

model6.add(Flatten())
model6.add(Dense(128, activation = 'relu'))
model6.add(Dense(256, activation = 'relu'))
model6.add(Dense(5, activation = 'softmax'))

model6.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model6.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_15 (MaxPooling)	(None, 40, 40, 16)	0
dropout_15 (Dropout)	(None, 40, 40, 16)	0
conv2d_16 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_16 (MaxPooling)	(None, 20, 20, 32)	0
dropout_16 (Dropout)	(None, 20, 20, 32)	0
conv2d_17 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_17 (MaxPooling)	(None, 10, 10, 64)	0
dropout_17 (Dropout)	(None, 10, 10, 64)	0
flatten_5 (Flatten)	(None, 6400)	0
dense_6 (Dense)	(None, 128)	819328
dense_7 (Dense)	(None, 256)	33024
dense_8 (Dense)	(None, 5)	1285

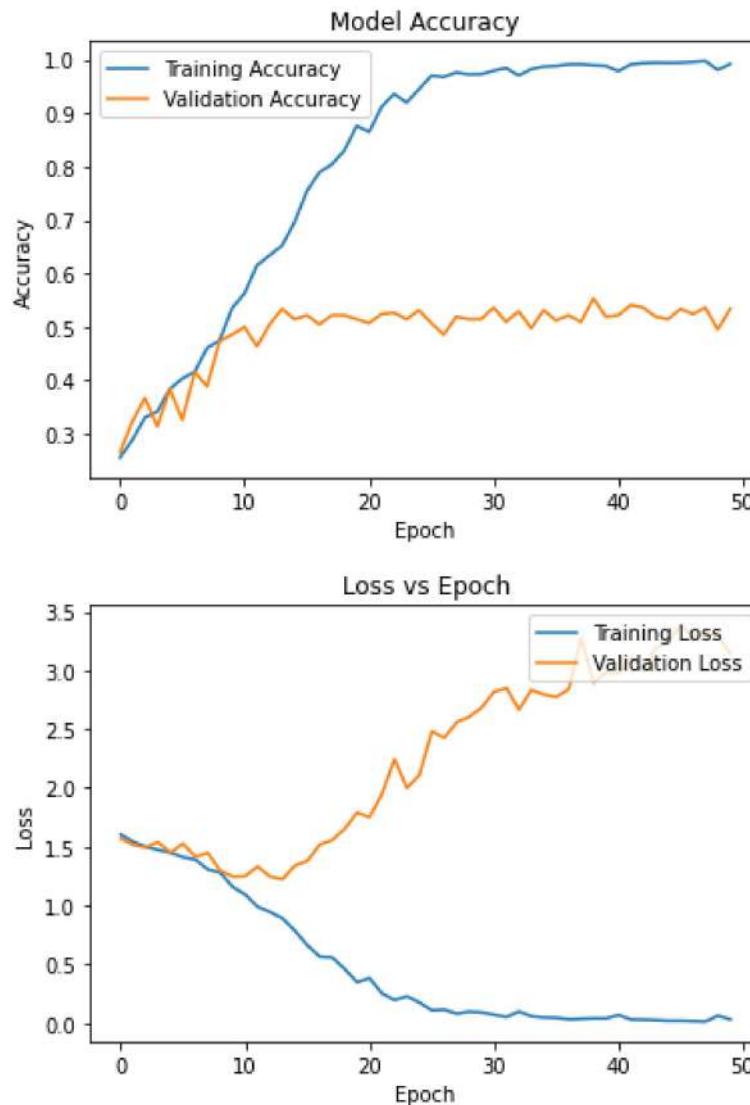
Total params: 917,893

Trainable params: 917,893

Non-trainable params: 0

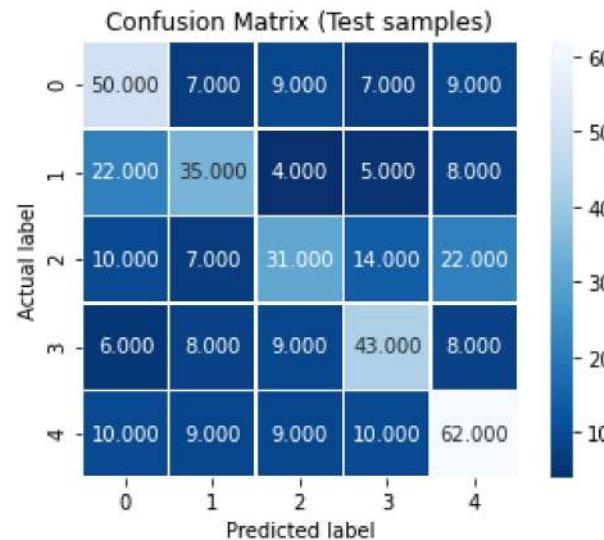
```
In [ ]: History6 = model6.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

```
In [ ]: draw(History6, model6, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.9975806474685669

Maximum validation accuracy: 0.5531401038169861



	precision	recall	f1-score	support
0	0.51	0.61	0.56	82
1	0.53	0.47	0.50	74
2	0.50	0.37	0.42	84
3	0.54	0.58	0.56	74
4	0.57	0.62	0.59	100
accuracy			0.53	414
macro avg	0.53	0.53	0.53	414
weighted avg	0.53	0.53	0.53	414

Model 6 is the best model with an accuracy of 58.82% on the validation dataset.

```
model6 = Sequential()

model6.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))

model6.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
```

```
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))
model6.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model6.add(MaxPooling2D(pool_size=(2, 2)))
model6.add(Dropout(0.1))

model6.add(Flatten())
model6.add(Dense(128, activation = 'relu'))
model6.add(Dense(256, activation = 'relu'))
model6.add(Dense(5, activation = 'softmax'))

model6.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model6.summary()
```

Now, we will be trying out average pooling on the same model instead of max pooling.

Model 7

```
In [ ]: model7 = Sequential()

model7.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='relu', input_shape = (80,80,1)))
model7.add(AveragePooling2D(pool_size=(2, 2)))
model7.add(Dropout(0.1))

model7.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model7.add(AveragePooling2D(pool_size=(2, 2)))
model7.add(Dropout(0.1))
model7.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='relu'))
model7.add(AveragePooling2D(pool_size=(2, 2)))
model7.add(Dropout(0.1))

model7.add(Flatten())
model7.add(Dense(128, activation = 'relu'))
model7.add(Dense(256, activation = 'relu'))
model7.add(Dense(5, activation = 'softmax'))
```

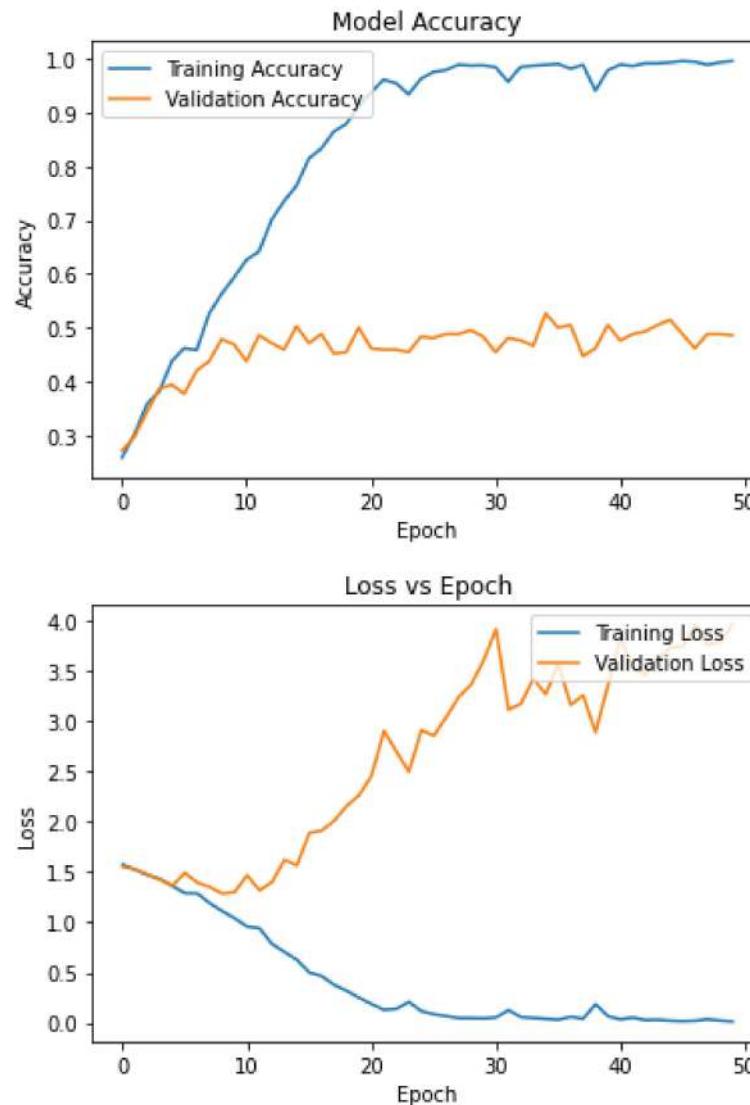
```
model7.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model7.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_18 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
average_pooling2d (AveragePo	(None, 40, 40, 16)	0
<hr/>		
dropout_18 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_19 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
average_pooling2d_1 (Average	(None, 20, 20, 32)	0
<hr/>		
dropout_19 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_20 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
average_pooling2d_2 (Average	(None, 10, 10, 64)	0
<hr/>		
dropout_20 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_6 (Flatten)	(None, 6400)	0
<hr/>		
dense_9 (Dense)	(None, 128)	819328
<hr/>		
dense_10 (Dense)	(None, 256)	33024
<hr/>		
dense_11 (Dense)	(None, 5)	1285
<hr/>		
Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		

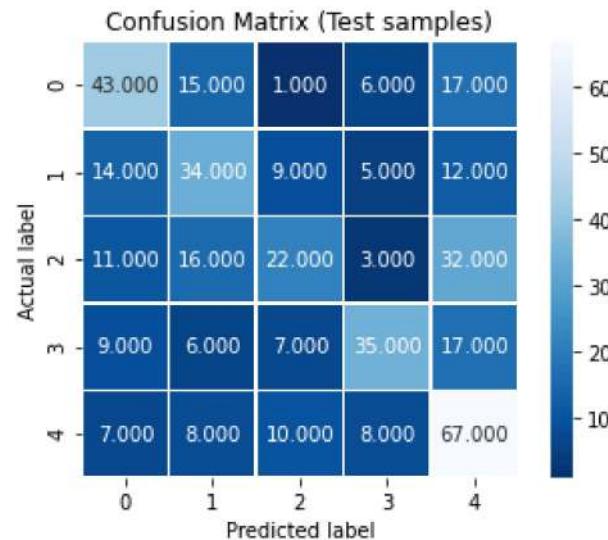
In []: History7 = model7.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History7, model7, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9962365627288818

Maximum validation accuracy: 0.5265700221061707



	precision	recall	f1-score	support
0	0.51	0.52	0.52	82
1	0.43	0.46	0.44	74
2	0.45	0.26	0.33	84
3	0.61	0.47	0.53	74
4	0.46	0.67	0.55	100
accuracy			0.49	414
macro avg	0.49	0.48	0.47	414
weighted avg	0.49	0.49	0.48	414

We see that average pooling does not give us better results than max pooling. Hence, we stick with model 6 and vary the activation functions to see which one performs better on model 6.

Model 8

In []: model8 = Sequential()

```
model8.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='sigmoid', input_shape = (80,80,1)))
model8.add(MaxPooling2D(pool_size=(2, 2)))
model8.add(Dropout(0.1))

model8.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='sigmoid'))
model8.add(MaxPooling2D(pool_size=(2, 2)))
model8.add(Dropout(0.1))
model8.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='sigmoid'))
model8.add(MaxPooling2D(pool_size=(2, 2)))
model8.add(Dropout(0.1))

model8.add(Flatten())
model8.add(Dense(128, activation ='sigmoid'))
model8.add(Dense(256, activation ='sigmoid'))
model8.add(Dense(5, activation = 'softmax'))

model8.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model8.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
conv2d_21 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_18 (MaxPooling)	(None, 40, 40, 16)	0
dropout_21 (Dropout)	(None, 40, 40, 16)	0
conv2d_22 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_19 (MaxPooling)	(None, 20, 20, 32)	0
dropout_22 (Dropout)	(None, 20, 20, 32)	0
conv2d_23 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_20 (MaxPooling)	(None, 10, 10, 64)	0
dropout_23 (Dropout)	(None, 10, 10, 64)	0
flatten_7 (Flatten)	(None, 6400)	0
dense_12 (Dense)	(None, 128)	819328
dense_13 (Dense)	(None, 256)	33024
dense_14 (Dense)	(None, 5)	1285

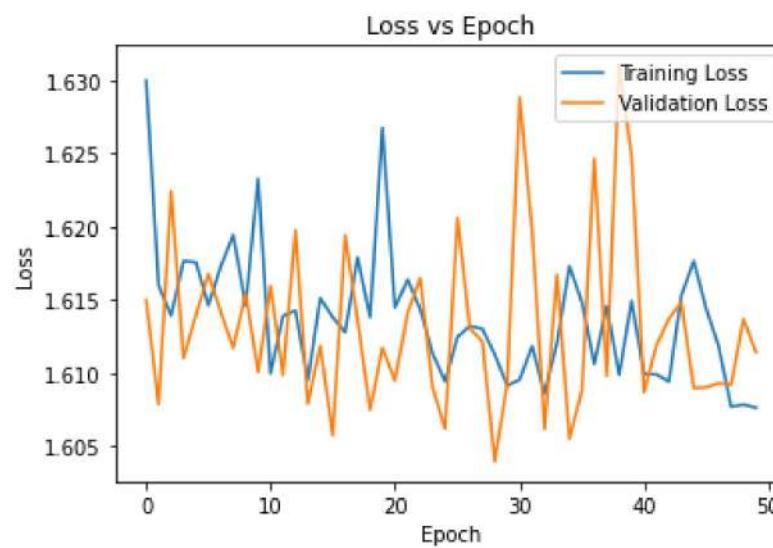
Total params: 917,893

Trainable params: 917,893

Non-trainable params: 0

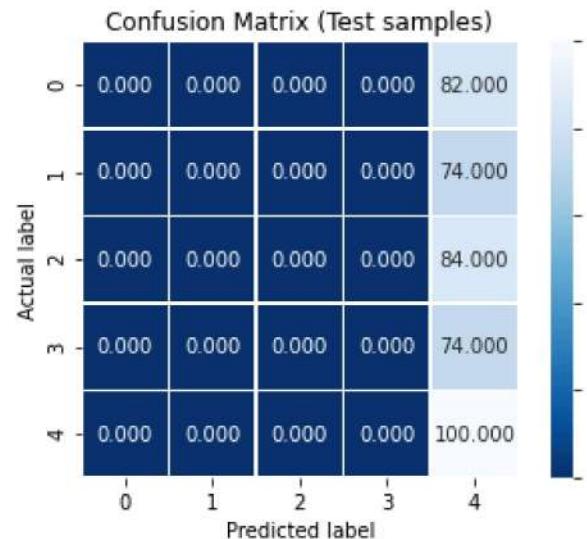
In []: History8 = model8.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History8, model8, train, test, train_labels, test_labels)



Maximum training accuracy: 0.23198924958705902

Maximum validation accuracy: 0.2415459007024765



	precision	recall	f1-score	support
0	0.00	0.00	0.00	82
1	0.00	0.00	0.00	74
2	0.00	0.00	0.00	84
3	0.00	0.00	0.00	74
4	0.24	1.00	0.39	100
accuracy			0.24	414
macro avg	0.05	0.20	0.08	414
weighted avg	0.06	0.24	0.09	414

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

Model 9

```
In [ ]: model9 = Sequential()

model9.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation ='elu', input_shape = (80,80,1)))
model9.add(MaxPooling2D(pool_size=(2, 2)))
model9.add(Dropout(0.1))
```

```
model9.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation ='elu'))
model9.add(MaxPooling2D(pool_size=(2, 2)))
model9.add(Dropout(0.1))
model9.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation ='elu'))
model9.add(MaxPooling2D(pool_size=(2, 2)))
model9.add(Dropout(0.1))

model9.add(Flatten())
model9.add(Dense(128, activation ='elu'))
model9.add(Dense(256, activation ='elu'))
model9.add(Dense(5, activation = 'softmax'))

model9.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model9.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_24 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
max_pooling2d_21 (MaxPooling)	(None, 40, 40, 16)	0
<hr/>		
dropout_24 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_25 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
max_pooling2d_22 (MaxPooling)	(None, 20, 20, 32)	0
<hr/>		
dropout_25 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_26 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_23 (MaxPooling)	(None, 10, 10, 64)	0
<hr/>		
dropout_26 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_8 (Flatten)	(None, 6400)	0
<hr/>		
dense_15 (Dense)	(None, 128)	819328
<hr/>		
dense_16 (Dense)	(None, 256)	33024
<hr/>		
dense_17 (Dense)	(None, 5)	1285
<hr/>		

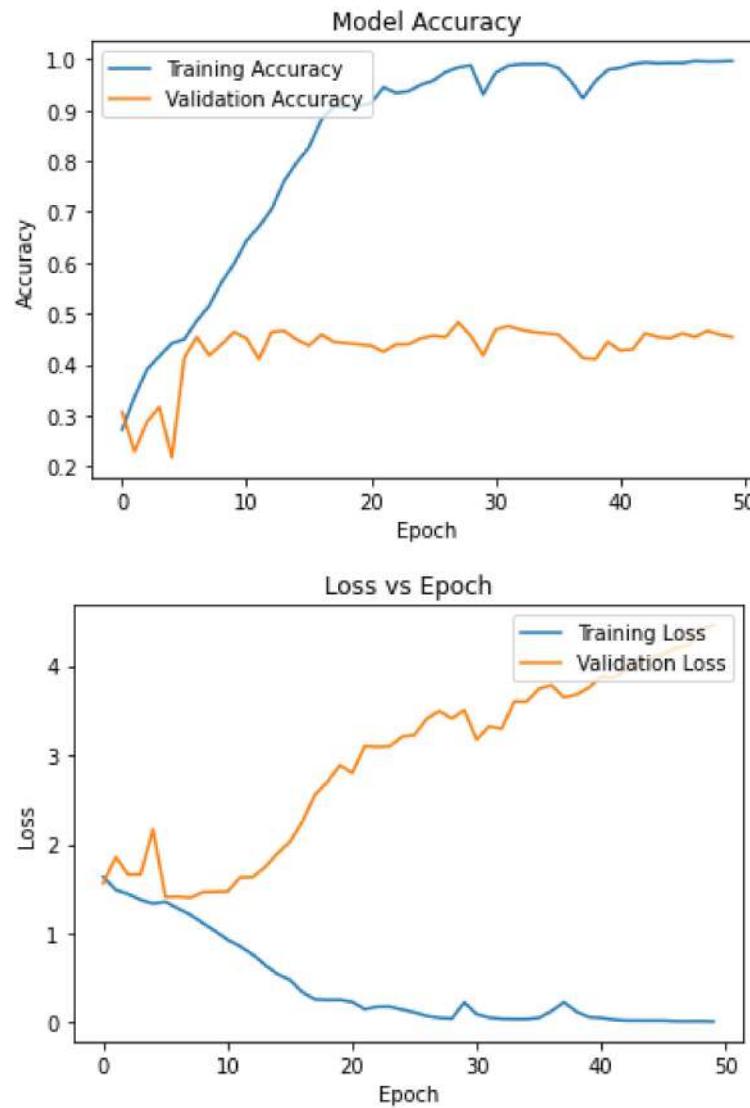
Total params: 917,893

Trainable params: 917,893

Non-trainable params: 0

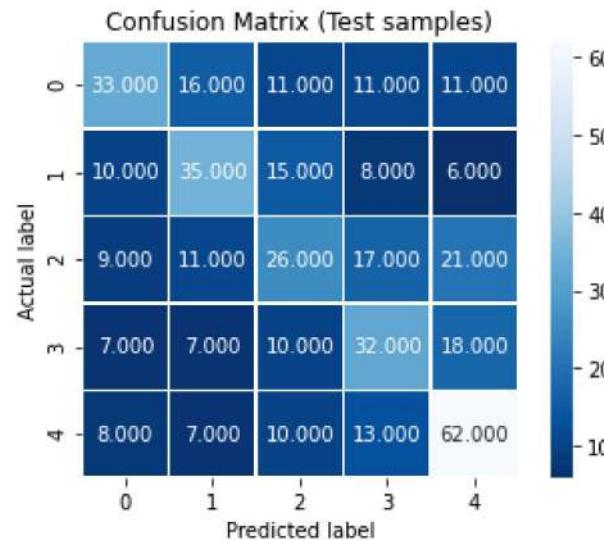
In []: History9 = model9.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History9, model9, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9962365627288818

Maximum validation accuracy: 0.483091801404953



	precision	recall	f1-score	support
0	0.49	0.40	0.44	82
1	0.46	0.47	0.47	74
2	0.36	0.31	0.33	84
3	0.40	0.43	0.41	74
4	0.53	0.62	0.57	100
accuracy			0.45	414
macro avg	0.45	0.45	0.44	414
weighted avg	0.45	0.45	0.45	414

Model 10

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model10 = Sequential()

model10.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))
model10.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
```

```
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Flatten())
model10.add(Dense(128, activation = leakyrelu))
model10.add(Dense(256, activation = leakyrelu))
model10.add(Dense(5, activation = 'softmax'))

model10.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model10.summary()
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_27 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_24 (MaxPooling)	(None, 40, 40, 16)	0
dropout_27 (Dropout)	(None, 40, 40, 16)	0
conv2d_28 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_25 (MaxPooling)	(None, 20, 20, 32)	0
dropout_28 (Dropout)	(None, 20, 20, 32)	0
conv2d_29 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_26 (MaxPooling)	(None, 10, 10, 64)	0
dropout_29 (Dropout)	(None, 10, 10, 64)	0
flatten_9 (Flatten)	(None, 6400)	0
dense_18 (Dense)	(None, 128)	819328
dense_19 (Dense)	(None, 256)	33024
dense_20 (Dense)	(None, 5)	1285

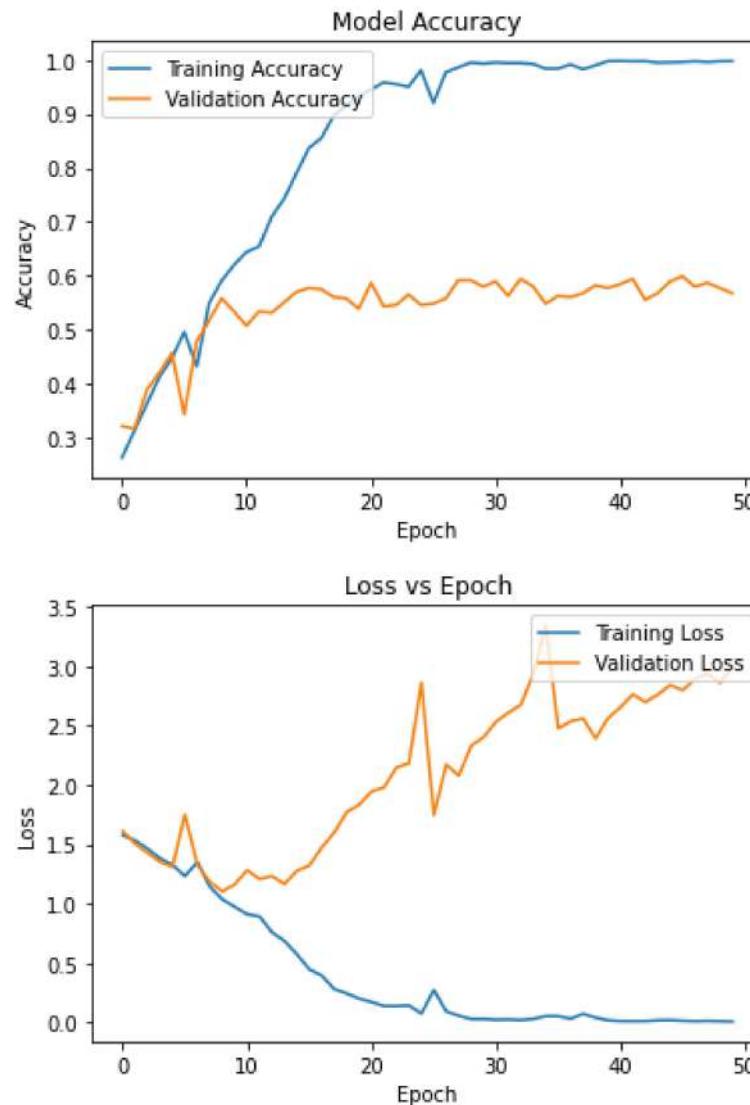
Total params: 917,893

Trainable params: 917,893

Non-trainable params: 0

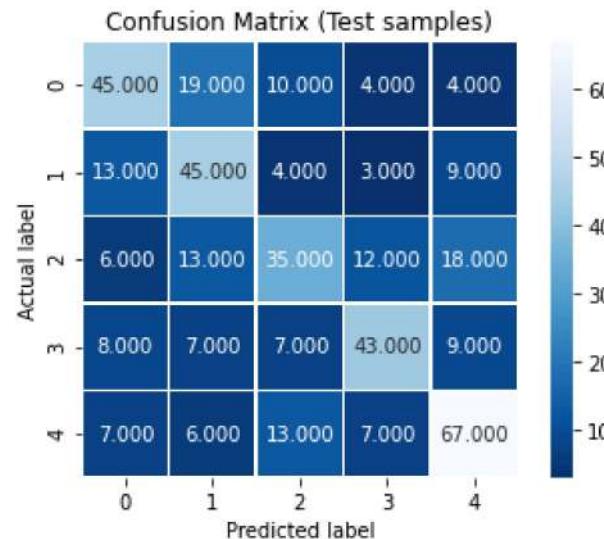
In []: History10 = model10.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History10, model10, train, test, train_labels, test_labels)



Maximum training accuracy: 0.9981182813644409

Maximum validation accuracy: 0.5990338325500488



	precision	recall	f1-score	support
0	0.57	0.55	0.56	82
1	0.50	0.61	0.55	74
2	0.51	0.42	0.46	84
3	0.62	0.58	0.60	74
4	0.63	0.67	0.65	100
accuracy			0.57	414
macro avg	0.57	0.56	0.56	414
weighted avg	0.57	0.57	0.57	414

After modifying the activation functions, we see that model 10's accuracy is the highest.

```
leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model10 = Sequential()

model10.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))
```

```
model10.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))
model10.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model10.add(MaxPooling2D(pool_size=(2, 2)))
model10.add(Dropout(0.1))

model10.add(Flatten())
model10.add(Dense(128, activation = leakyrelu))
model10.add(Dense(256, activation = leakyrelu))
model10.add(Dense(5, activation = 'softmax'))

model10.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model10.summary()
```

Thus, we will now vary the regularization parameters for model 10.

Model 11

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model11 = Sequential()

model11.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model11.add(MaxPooling2D(pool_size=(2, 2)))
model11.add(Dropout(0.25))

model11.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model11.add(MaxPooling2D(pool_size=(2, 2)))
model11.add(Dropout(0.25))

model11.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model11.add(MaxPooling2D(pool_size=(2, 2)))
model11.add(Dropout(0.25))

model11.add(Flatten())
model11.add(Dense(128, activation = leakyrelu))
model11.add(Dense(256, activation = leakyrelu))
```

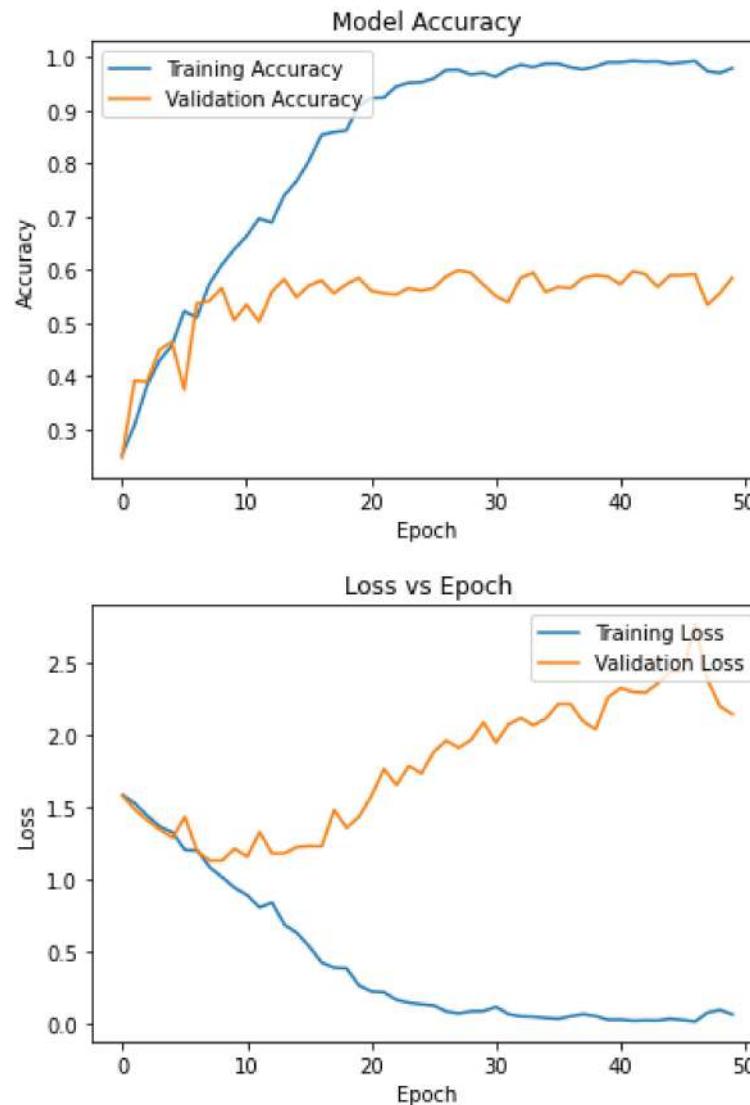
```
model11.add(Dense(5, activation = 'softmax'))  
  
model11.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])  
model11.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_30 (Conv2D)	(None, 80, 80, 16)	160
<hr/>		
max_pooling2d_27 (MaxPooling)	(None, 40, 40, 16)	0
<hr/>		
dropout_30 (Dropout)	(None, 40, 40, 16)	0
<hr/>		
conv2d_31 (Conv2D)	(None, 40, 40, 32)	12832
<hr/>		
max_pooling2d_28 (MaxPooling)	(None, 20, 20, 32)	0
<hr/>		
dropout_31 (Dropout)	(None, 20, 20, 32)	0
<hr/>		
conv2d_32 (Conv2D)	(None, 20, 20, 64)	51264
<hr/>		
max_pooling2d_29 (MaxPooling)	(None, 10, 10, 64)	0
<hr/>		
dropout_32 (Dropout)	(None, 10, 10, 64)	0
<hr/>		
flatten_10 (Flatten)	(None, 6400)	0
<hr/>		
dense_21 (Dense)	(None, 128)	819328
<hr/>		
dense_22 (Dense)	(None, 256)	33024
<hr/>		
dense_23 (Dense)	(None, 5)	1285
<hr/>		
Total params: 917,893		
Trainable params: 917,893		
Non-trainable params: 0		

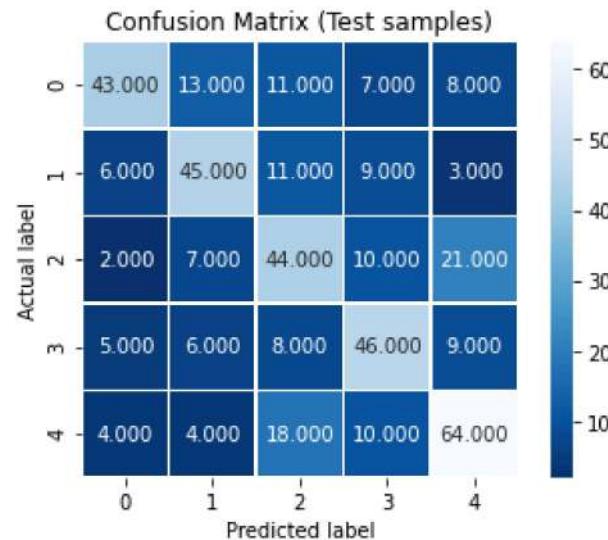
```
In [ ]: History11 = model11.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

```
In [ ]: draw(History11, model11, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.9927419424057007

Maximum validation accuracy: 0.5990338325500488



	precision	recall	f1-score	support
0	0.72	0.52	0.61	82
1	0.60	0.61	0.60	74
2	0.48	0.52	0.50	84
3	0.56	0.62	0.59	74
4	0.61	0.64	0.62	100
accuracy			0.58	414
macro avg	0.59	0.58	0.58	414
weighted avg	0.59	0.58	0.59	414

Model 12

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model12 = Sequential()

model12.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model12.add(MaxPooling2D(pool_size=(2, 2)))

model12.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model12.add(MaxPooling2D(pool_size=(2, 2)))
model12.add(BatchNormalization())

model12.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
```

```
model12.add(MaxPooling2D(pool_size=(2, 2)))
model12.add(BatchNormalization())

model12.add(Flatten())
model12.add(Dense(128, activation = leakyrelu))
model12.add(Dense(256, activation = leakyrelu))
model12.add(Dense(5, activation = 'softmax'))

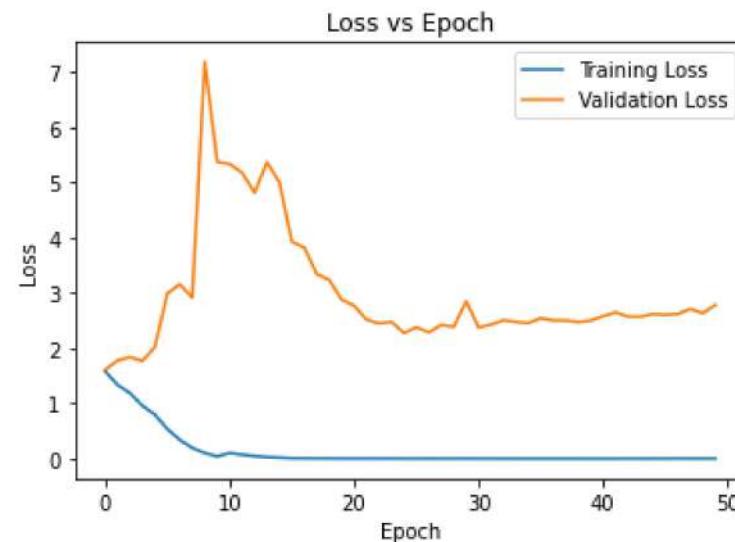
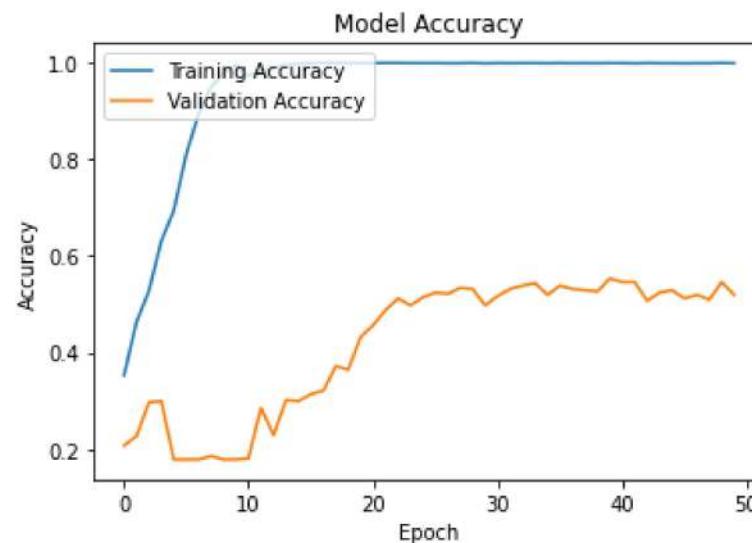
model12.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
model12.summary()
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv2d_33 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_30 (MaxPooling)	(None, 40, 40, 16)	0
conv2d_34 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_31 (MaxPooling)	(None, 20, 20, 32)	0
batch_normalization (BatchNo)	(None, 20, 20, 32)	128
conv2d_35 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_32 (MaxPooling)	(None, 10, 10, 64)	0
batch_normalization_1 (Batch)	(None, 10, 10, 64)	256
flatten_11 (Flatten)	(None, 6400)	0
dense_24 (Dense)	(None, 128)	819328
dense_25 (Dense)	(None, 256)	33024
dense_26 (Dense)	(None, 5)	1285
<hr/>		
Total params: 918,277		
Trainable params: 918,085		
Non-trainable params: 192		

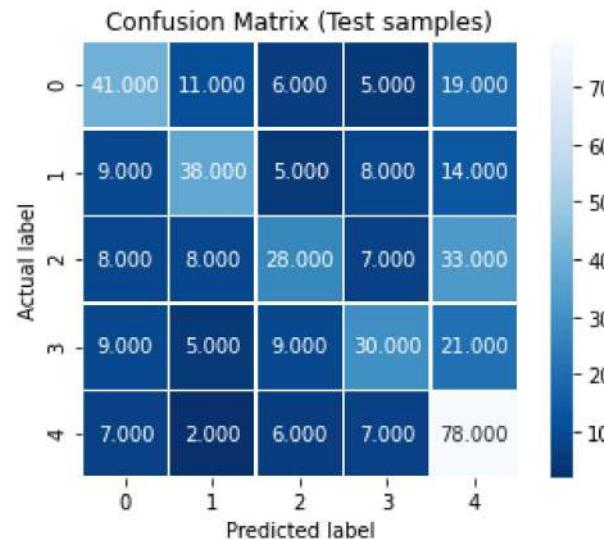
```
In [ ]: History12 = model12.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

```
In [ ]: draw(History12, model12, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.998924732208252

Maximum validation accuracy: 0.5531401038169861



	precision	recall	f1-score	support
0	0.55	0.50	0.53	82
1	0.59	0.51	0.55	74
2	0.52	0.33	0.41	84
3	0.53	0.41	0.46	74
4	0.47	0.78	0.59	100
accuracy			0.52	414
macro avg	0.53	0.51	0.51	414
weighted avg	0.53	0.52	0.51	414

Model 13

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model13 = Sequential()

model13.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model13.add(MaxPooling2D(pool_size=(2, 2)))
model13.add(Dropout(0.1))

model13.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model13.add(MaxPooling2D(pool_size=(2, 2)))
model13.add(BatchNormalization())
model13.add(Dropout(0.1))
```

```
model13.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model13.add(MaxPooling2D(pool_size=(2, 2)))
model13.add(BatchNormalization())
model13.add(Dropout(0.1))

model13.add(Flatten())
model13.add(Dense(128, activation = leakyrelu))
model13.add(Dense(256, activation = leakyrelu))
model13.add(Dense(5, activation = 'softmax'))

model13.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model13.summary()
```

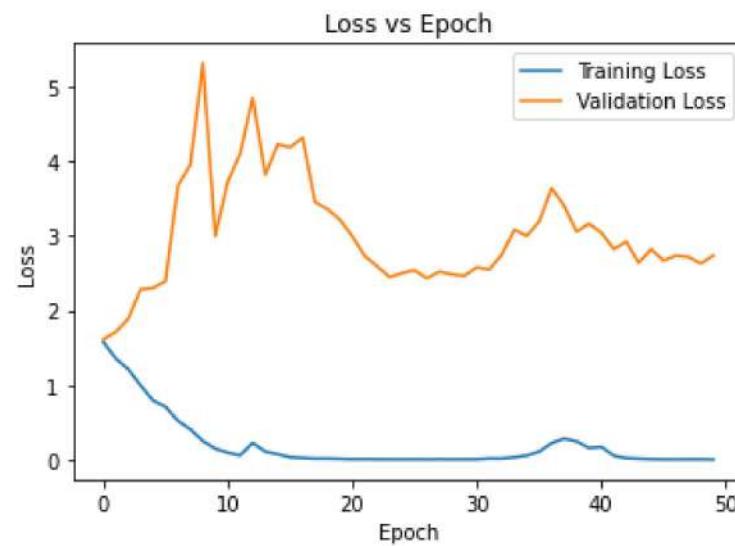
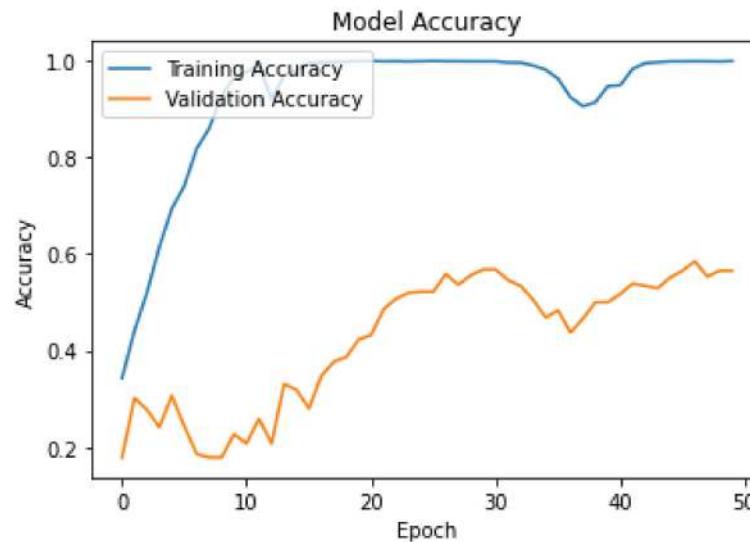
Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_36 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_33 (MaxPooling)	(None, 40, 40, 16)	0
dropout_33 (Dropout)	(None, 40, 40, 16)	0
conv2d_37 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_34 (MaxPooling)	(None, 20, 20, 32)	0
batch_normalization_2 (Batch)	(None, 20, 20, 32)	128
dropout_34 (Dropout)	(None, 20, 20, 32)	0
conv2d_38 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_35 (MaxPooling)	(None, 10, 10, 64)	0
batch_normalization_3 (Batch)	(None, 10, 10, 64)	256
dropout_35 (Dropout)	(None, 10, 10, 64)	0
flatten_12 (Flatten)	(None, 6400)	0
dense_27 (Dense)	(None, 128)	819328
dense_28 (Dense)	(None, 256)	33024
dense_29 (Dense)	(None, 5)	1285

Total params: 918,277
Trainable params: 918,085
Non-trainable params: 192

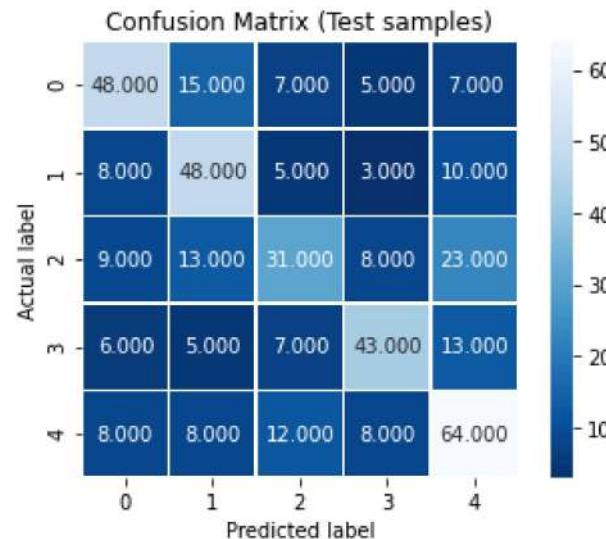
In []: History13 = model13.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

In []: draw(History13, model13, train, test, train_labels, test_labels)



Maximum training accuracy: 0.998924732208252

Maximum validation accuracy: 0.5845410823822021



	precision	recall	f1-score	support
0	0.61	0.59	0.60	82
1	0.54	0.65	0.59	74
2	0.50	0.37	0.42	84
3	0.64	0.58	0.61	74
4	0.55	0.64	0.59	100
accuracy			0.57	414
macro avg	0.57	0.56	0.56	414
weighted avg	0.57	0.57	0.56	414

Model 10 still performs the best. Now, we will be adding more convolution layers to it and check if it performs better.

Model 14

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model14 = Sequential()

model14.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model14.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model14.add(Dropout(0.1))

model14.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model14.add(MaxPooling2D(pool_size=(2, 2)))
model14.add(Dropout(0.1))

model14.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model14.add(MaxPooling2D(pool_size=(2, 2)))
model14.add(Dropout(0.1))

model14.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model14.add(MaxPooling2D(pool_size=(2, 2)))
model14.add(Dropout(0.1))

model14.add(Flatten())
model14.add(Dense(128, activation = leakyrelu))
model14.add(Dense(256, activation = leakyrelu))
model14.add(Dense(5, activation = 'softmax'))

model14.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model14.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
conv2d_39 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_36 (MaxPooling)	(None, 40, 40, 16)	0
dropout_36 (Dropout)	(None, 40, 40, 16)	0
conv2d_40 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_37 (MaxPooling)	(None, 20, 20, 32)	0
dropout_37 (Dropout)	(None, 20, 20, 32)	0
conv2d_41 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_38 (MaxPooling)	(None, 10, 10, 64)	0
dropout_38 (Dropout)	(None, 10, 10, 64)	0
conv2d_42 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_39 (MaxPooling)	(None, 5, 5, 128)	0
dropout_39 (Dropout)	(None, 5, 5, 128)	0
flatten_13 (Flatten)	(None, 3200)	0
dense_30 (Dense)	(None, 128)	409728
dense_31 (Dense)	(None, 256)	33024
dense_32 (Dense)	(None, 5)	1285

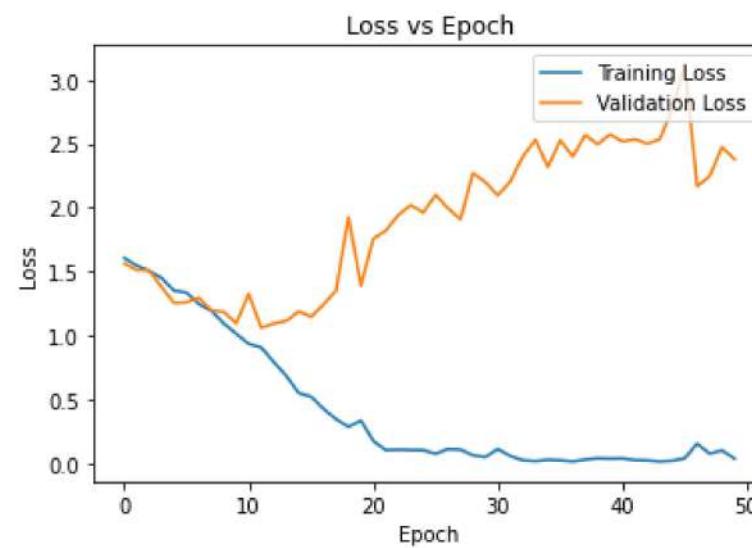
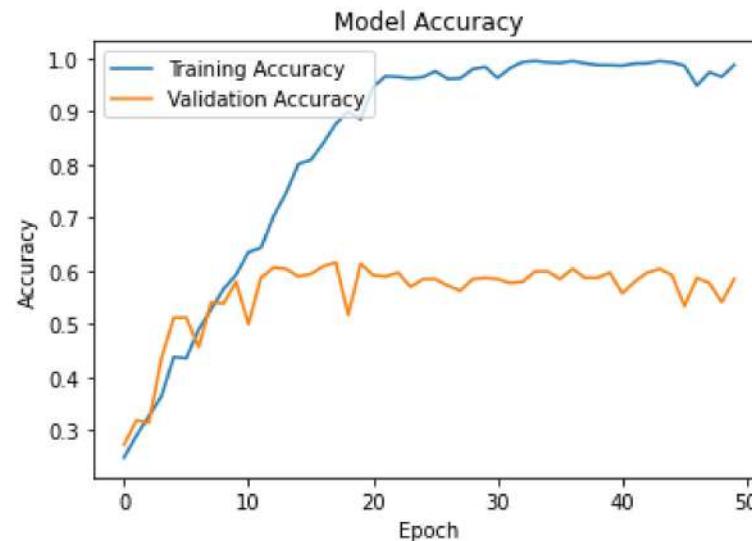
Total params: 713,221

Trainable params: 713,221

Non-trainable params: 0

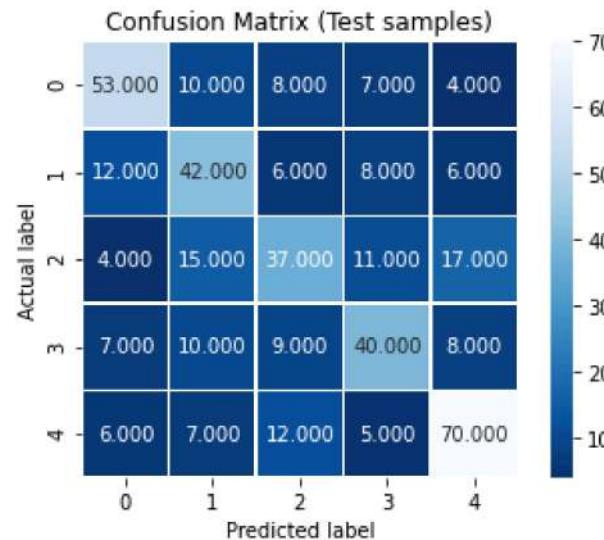
In []: History14 = model14.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))

```
In [ ]: draw(History14, model14, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.9951612949371338

Maximum validation accuracy: 0.6159420013427734



	precision	recall	f1-score	support
0	0.65	0.65	0.65	82
1	0.50	0.57	0.53	74
2	0.51	0.44	0.47	84
3	0.56	0.54	0.55	74
4	0.67	0.70	0.68	100
accuracy			0.58	414
macro avg	0.58	0.58	0.58	414
weighted avg	0.58	0.58	0.58	414

Model 15

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model15 = Sequential()

model15.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model15.add(MaxPooling2D(pool_size=(2, 2)))
model15.add(Dropout(0.1))

model15.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model15.add(MaxPooling2D(pool_size=(2, 2)))
model15.add(Dropout(0.1))
```

```
model15.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model15.add(MaxPooling2D(pool_size=(2, 2)))
model15.add(Dropout(0.1))

model15.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model15.add(MaxPooling2D(pool_size=(2, 2)))
model15.add(Dropout(0.1))

model15.add(Conv2D(filters = 256, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model15.add(MaxPooling2D(pool_size=(2, 2)))
model15.add(Dropout(0.1))

model15.add(Flatten())
model15.add(Dense(128, activation = leakyrelu))
model15.add(Dense(256, activation = leakyrelu))
model15.add(Dense(5, activation = 'softmax'))

model15.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model15.summary()
```

Model: "sequential_14"

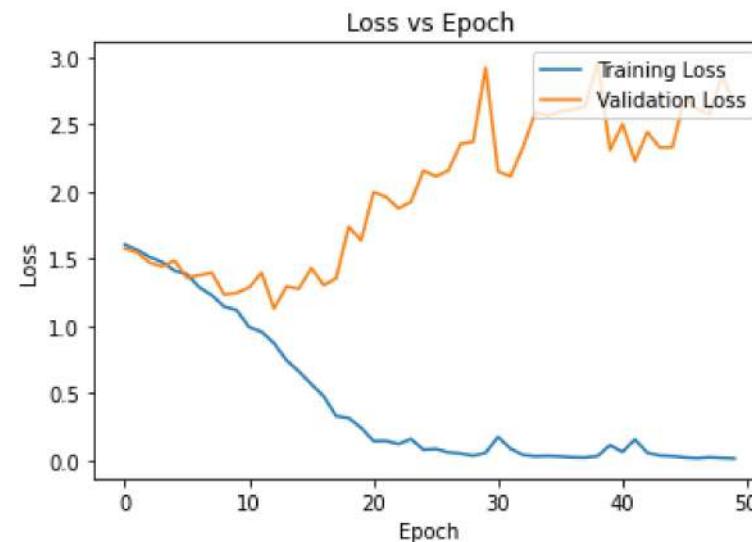
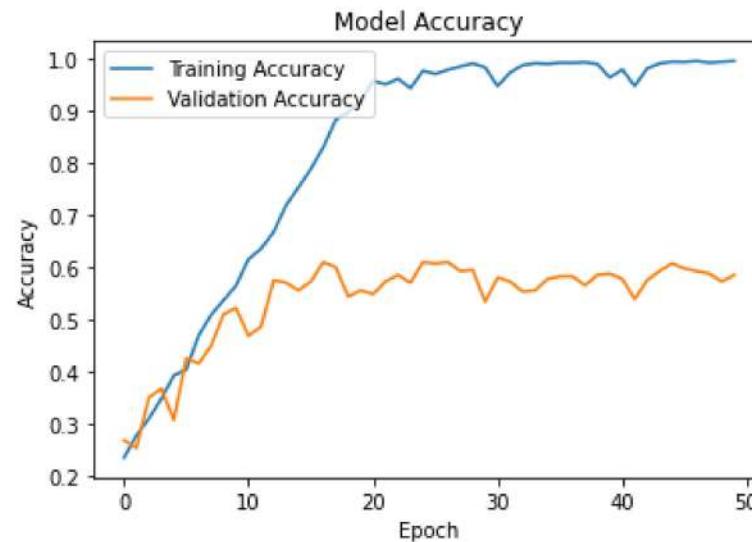
Layer (type)	Output Shape	Param #
conv2d_43 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_40 (MaxPooling)	(None, 40, 40, 16)	0
dropout_40 (Dropout)	(None, 40, 40, 16)	0
conv2d_44 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_41 (MaxPooling)	(None, 20, 20, 32)	0
dropout_41 (Dropout)	(None, 20, 20, 32)	0
conv2d_45 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_42 (MaxPooling)	(None, 10, 10, 64)	0
dropout_42 (Dropout)	(None, 10, 10, 64)	0
conv2d_46 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_43 (MaxPooling)	(None, 5, 5, 128)	0
dropout_43 (Dropout)	(None, 5, 5, 128)	0
conv2d_47 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_44 (MaxPooling)	(None, 2, 2, 256)	0
dropout_44 (Dropout)	(None, 2, 2, 256)	0
flatten_14 (Flatten)	(None, 1024)	0
dense_33 (Dense)	(None, 128)	131200
dense_34 (Dense)	(None, 256)	33024
dense_35 (Dense)	(None, 5)	1285
<hr/>		
Total params: 1,254,149		

Trainable params: 1,254,149

Non-trainable params: 0

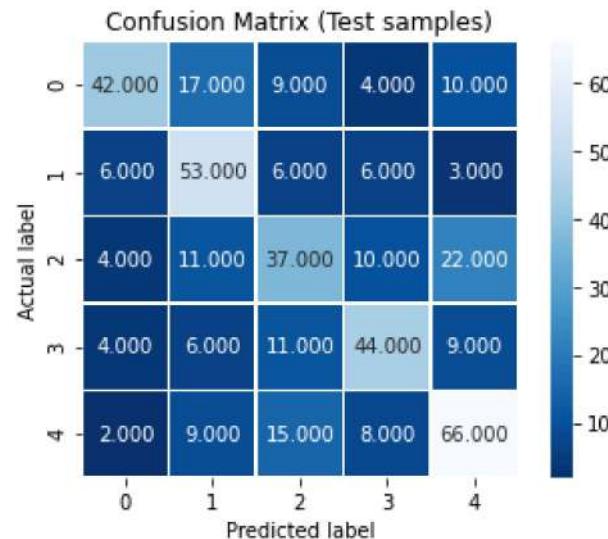
```
In [ ]: History15 = model15.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

```
In [ ]: draw(History15, model15, train, test, train_labels, test_labels)
```



Maximum training accuracy: 0.9943548440933228

Maximum validation accuracy: 0.6086956262588501



	precision	recall	f1-score	support
0	0.72	0.51	0.60	82
1	0.55	0.72	0.62	74
2	0.47	0.44	0.46	84
3	0.61	0.59	0.60	74
4	0.60	0.66	0.63	100
accuracy			0.58	414
macro avg	0.59	0.58	0.58	414
weighted avg	0.59	0.58	0.58	414

Model 16

```
In [26]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model16 = Sequential()

model16.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))
```

```
model16.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 256, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 512, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Flatten())
model16.add(Dense(128, activation = leakyrelu))
model16.add(Dense(256, activation = leakyrelu))
model16.add(Dense(5, activation = 'softmax'))

model16.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=[ 'accuracy'])
model16.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 80, 80, 16)	160
max_pooling2d_3 (MaxPooling 2D)	(None, 40, 40, 16)	0
dropout_3 (Dropout)	(None, 40, 40, 16)	0
conv2d_4 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_4 (MaxPooling 2D)	(None, 20, 20, 32)	0
dropout_4 (Dropout)	(None, 20, 20, 32)	0
conv2d_5 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_5 (MaxPooling 2D)	(None, 10, 10, 64)	0
dropout_5 (Dropout)	(None, 10, 10, 64)	0
conv2d_6 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_6 (MaxPooling 2D)	(None, 5, 5, 128)	0
dropout_6 (Dropout)	(None, 5, 5, 128)	0
conv2d_7 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_7 (MaxPooling 2D)	(None, 2, 2, 256)	0
dropout_7 (Dropout)	(None, 2, 2, 256)	0
conv2d_8 (Conv2D)	(None, 2, 2, 512)	3277312
max_pooling2d_8 (MaxPooling 2D)	(None, 1, 1, 512)	0

```
dropout_8 (Dropout)      (None, 1, 1, 512)      0
flatten_1 (Flatten)      (None, 512)            0
dense_1 (Dense)          (None, 128)             65664
dense_2 (Dense)          (None, 256)             33024
dense_3 (Dense)          (None, 5)               1285
```

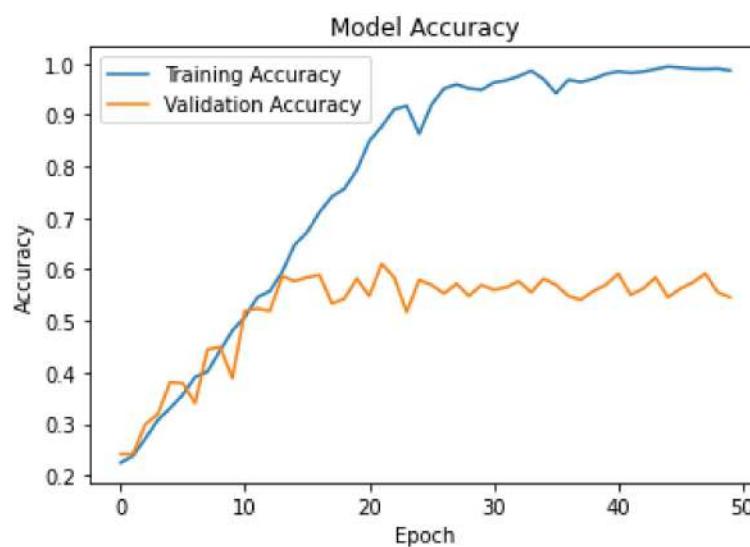
```
=====
Total params: 4,465,925
```

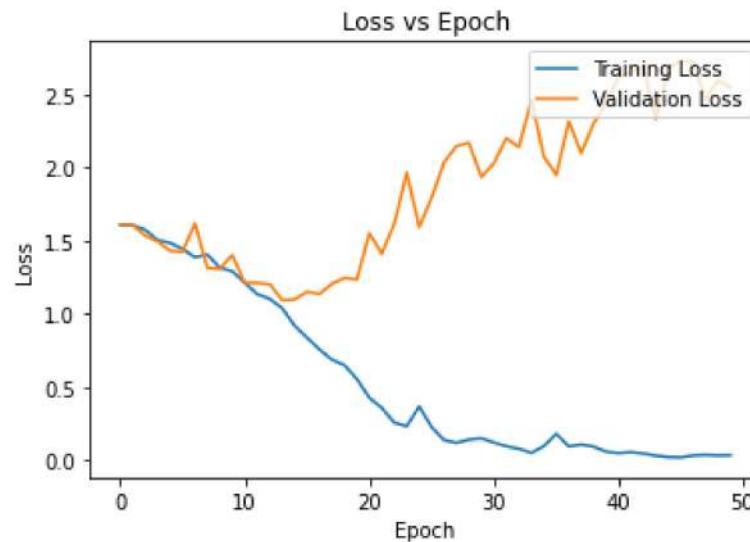
```
Trainable params: 4,465,925
```

```
Non-trainable params: 0
```

```
In [ ]: History16 = model16.fit(train,train_labels, batch_size=128, epochs = 50, validation_data = (test,test_labels))
```

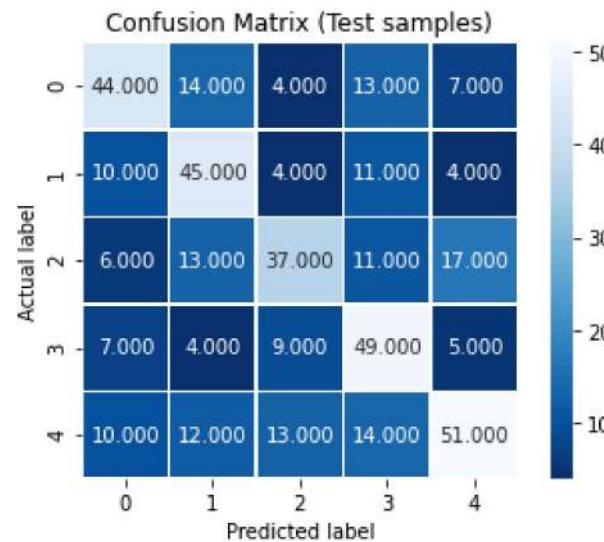
```
In [ ]: draw(History16, model16, train, test, train_labels, test_labels)
```





Maximum training accuracy: 0.9943548440933228

Maximum validation accuracy: 0.6111111044883728



	precision	recall	f1-score	support
0	0.57	0.54	0.55	82
1	0.51	0.61	0.56	74
2	0.55	0.44	0.49	84
3	0.50	0.66	0.57	74
4	0.61	0.51	0.55	100
accuracy			0.55	414
macro avg	0.55	0.55	0.54	414
weighted avg	0.55	0.55	0.54	414

Model 16 is hence the best model so far with an accuracy of 63.87% on the validation dataset.

```

leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model16 = Sequential()

model16.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,1)))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Conv2D(filters = 256, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

```

```
model16.add(Conv2D(filters = 512, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model16.add(MaxPooling2D(pool_size=(2, 2)))
model16.add(Dropout(0.1))

model16.add(Flatten())
model16.add(Dense(128, activation = leakyrelu))
model16.add(Dense(256, activation = leakyrelu))
model16.add(Dense(5, activation = 'softmax'))

model16.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model16.summary()
```

Now, we will use this model and run the experiment on color images.

Load the colour dataset

```
In [29]: X = []
Z = []
IMG_SIZE = 80
BASE_PATH = "/content/drive/My Drive/ML_DRIVE/Assign_6/flowers_dataset/"
FLOWER_DAISY_DIR      = BASE_PATH + "daisy"
FLOWER_SUNFLOWER_DIR   = BASE_PATH + "sunflower"
FLOWER_TULIP_DIR       = BASE_PATH + "tulip"
FLOWER_DANDI_DIR       = BASE_PATH + "dandelion"
FLOWER_ROSE_DIR        = BASE_PATH + "rose"
```

```
In [ ]: make_train_data('Daisy',FLOWER_DAISY_DIR, 'color')
```

```
In [ ]: make_train_data('Sunflower',FLOWER_SUNFLOWER_DIR, 'color')
```

```
In [ ]: make_train_data('Tulip',FLOWER_TULIP_DIR, 'color')
```

```
In [ ]: make_train_data('Dandelion',FLOWER_DANDI_DIR, 'color')
```

```
In [ ]: make_train_data('Rose',FLOWER_ROSE_DIR, 'color')
```

```
In [ ]: def draw_matrix(model, train, test, train_labels, test_labels):
    predictions_test = []
    for item in test:
        predictions_test.append(np.argmax(model.predict(item.reshape(-1,80,80,3)))))

    cm_test = sklearn.metrics.confusion_matrix(np.argmax(test_labels, axis = 1), predictions_test)
    sns.heatmap(cm_test, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues_r');
    plt.ylabel('Actual label');
    plt.xlabel('Predicted label');
    all_sample_title = 'Confusion Matrix (Test samples)'
    plt.title(all_sample_title)
    plt.show()
    print(classification_report(np.argmax(test_labels, axis = 1), predictions_test))

def draw(History, model, train, test, train_labels, test_labels):
    draw_graph(History)
    draw_matrix(model, train, test, train_labels, test_labels)
```

Pre-processing

```
In [ ]: le = LabelEncoder()
Y = le.fit_transform(Z)
Y = to_categorical(Y,5)
X = np.array(X)
X = X/255
```

Train : Test split

```
In [ ]: train, test, train_labels, test_labels = train_test_split(X, Y, test_size=0.1, random_state=100)
```

Training the model

the model is same as model 16(best) but in that input_shape was (80,80,1) as it was not colour. But since we require colour (80,80,3), hence we are creating the same model with change only in input_shape

```
In [ ]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model17 = Sequential()
```

```
model17.add(Conv2D(filters = 16, kernel_size = (3,3),padding = 'Same',activation = leakyrelu, input_shape = (80,80,3)))
model17.add(MaxPooling2D(pool_size=(2, 2)))
model17.add(Dropout(0.1))

model17.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model17.add(MaxPooling2D(pool_size=(2, 2)))
model17.add(Dropout(0.1))

model17.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model17.add(MaxPooling2D(pool_size=(2, 2)))
model17.add(Dropout(0.1))

model17.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model17.add(MaxPooling2D(pool_size=(2, 2)))
model17.add(Dropout(0.1))

model17.add(Conv2D(filters = 256, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model17.add(MaxPooling2D(pool_size=(2, 2)))
model17.add(Dropout(0.1))

model17.add(Flatten())
model17.add(Dense(128, activation = leakyrelu))
model17.add(Dense(256, activation = leakyrelu))
model17.add(Dense(5, activation = 'softmax'))

model17.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=['accuracy'])
model17.summary()
```

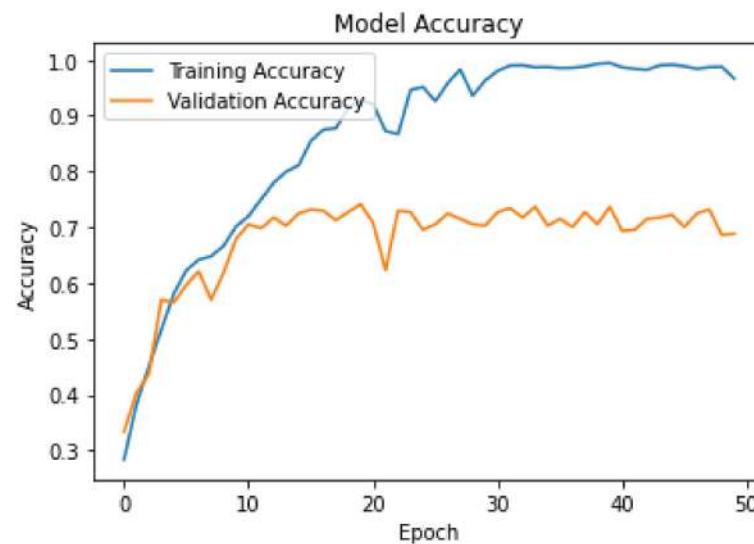
Model: "sequential_16"

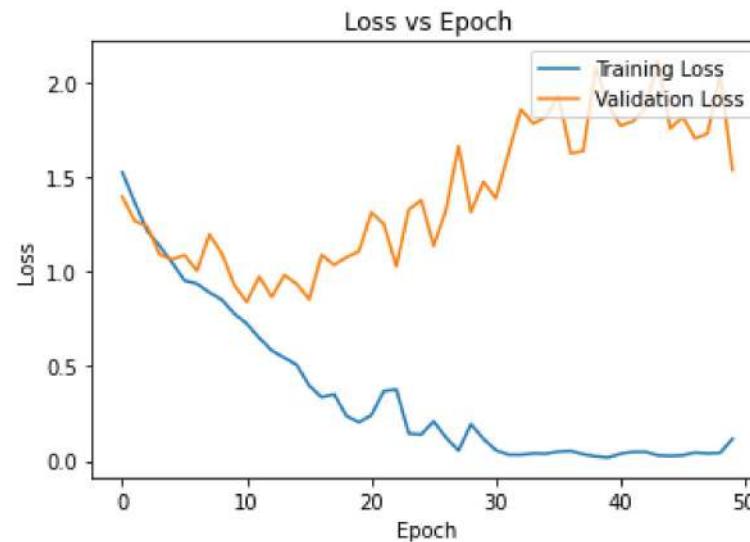
Layer (type)	Output Shape	Param #
conv2d_54 (Conv2D)	(None, 80, 80, 16)	448
max_pooling2d_51 (MaxPooling)	(None, 40, 40, 16)	0
dropout_51 (Dropout)	(None, 40, 40, 16)	0
conv2d_55 (Conv2D)	(None, 40, 40, 32)	12832
max_pooling2d_52 (MaxPooling)	(None, 20, 20, 32)	0
dropout_52 (Dropout)	(None, 20, 20, 32)	0
conv2d_56 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_53 (MaxPooling)	(None, 10, 10, 64)	0
dropout_53 (Dropout)	(None, 10, 10, 64)	0
conv2d_57 (Conv2D)	(None, 10, 10, 128)	204928
max_pooling2d_54 (MaxPooling)	(None, 5, 5, 128)	0
dropout_54 (Dropout)	(None, 5, 5, 128)	0
conv2d_58 (Conv2D)	(None, 5, 5, 256)	819456
max_pooling2d_55 (MaxPooling)	(None, 2, 2, 256)	0
dropout_55 (Dropout)	(None, 2, 2, 256)	0
conv2d_59 (Conv2D)	(None, 2, 2, 512)	3277312
max_pooling2d_56 (MaxPooling)	(None, 1, 1, 512)	0
dropout_56 (Dropout)	(None, 1, 1, 512)	0
flatten_16 (Flatten)	(None, 512)	0
dense_39 (Dense)	(None, 128)	65664

dense_40 (Dense)	(None, 256)	33024
dense_41 (Dense)	(None, 5)	1285
=====		
Total params: 4,466,213		
Trainable params: 4,466,213		
Non-trainable params: 0		

```
In [ ]: History17 = model17.fit(train, train_labels, batch_size=128, epochs=50, validation_data=(test,test_labels))
```

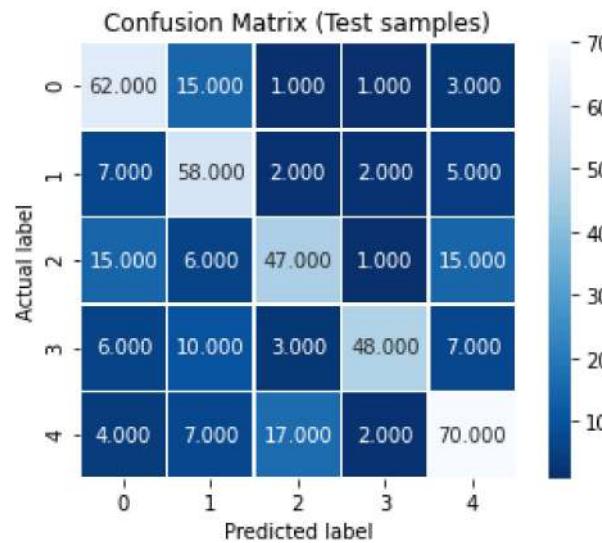
```
In [ ]: draw(History17, model17, train, test, train_labels, test_labels)
```





Maximum training accuracy: 0.9943548440933228

Maximum validation accuracy: 0.7415459156036377

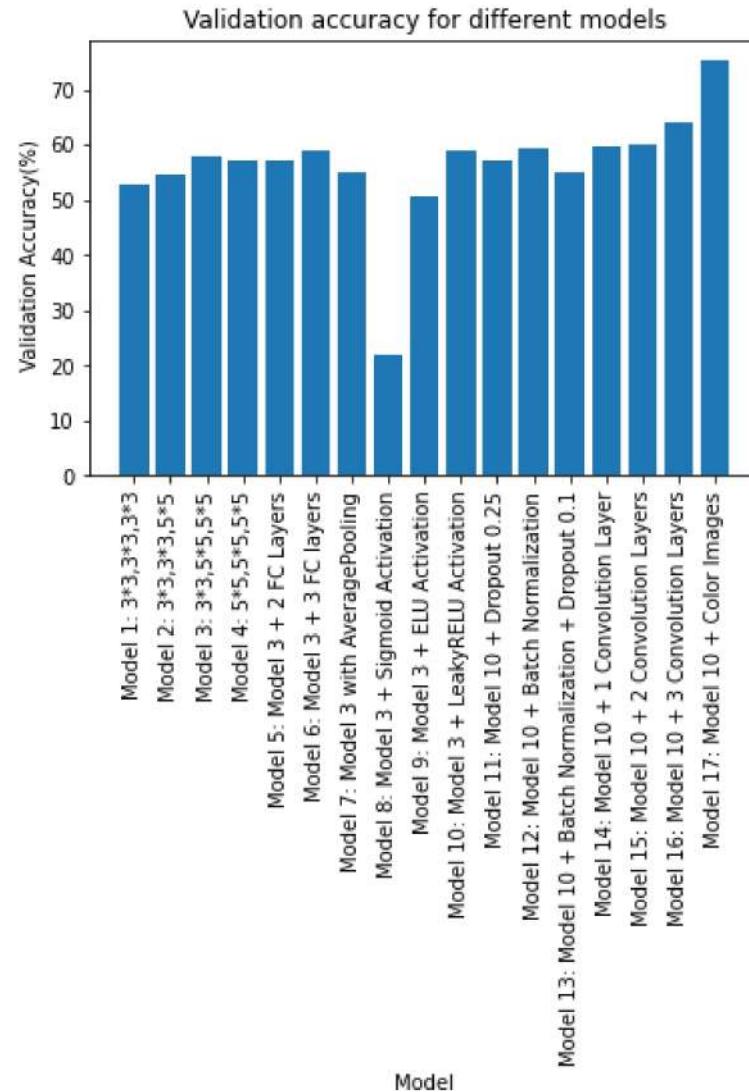


	precision	recall	f1-score	support
0	0.66	0.76	0.70	82
1	0.60	0.78	0.68	74
2	0.67	0.56	0.61	84
3	0.89	0.65	0.75	74
4	0.70	0.70	0.70	100
accuracy			0.69	414
macro avg	0.70	0.69	0.69	414
weighted avg	0.70	0.69	0.69	414

Plot bar-graph for comparing validation accuracy for all the models

```
In [ ]: validation = [52.94, 54.66, 57.84, 57.10, 57.11, 58.82, 55.14, 22.06, 50.49, 58.81, 57.11, 59.31, 54.90, 59.80, 60.05, 0
serial = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17]
labels = ['Model 1: 3*3,3*3,3*3', 'Model 2: 3*3,3*3,5*5', 'Model 3: 3*3,5*5,5*5', 'Model 4: 5*5,5*5,5*5', 'Model 5: Model 1 + 1 FC layer', 'Model 6: Model 3 + 3 FC layers', 'Model 7: Model 3 with AveragePooling', 'Model 8: Model 3 + Sigmoid Activation', 'Model 9: Model 3 + ELU Activation', 'Model 10: Model 3 + LeakyRELU Activation', 'Model 11: Model 10 + Dropout 0.1', 'Model 12: Model 10 + Batch Normalization', 'Model 13: Model 10 + Batch Normalization + Dropout 0.1', 'Model 14: Model 10 + 1 Convolution Layer', 'Model 15: Model 10 + 2 Convolution Layers', 'Model 16: Model 10 + 3 Convolution Layers', 'Model 17: Model 10 + Color Images']

plt.bar(serial,validation)
plt.title("Validation accuracy for different models")
plt.xticks(serial, labels, rotation = 90)
plt.xlabel("Model")
plt.ylabel(" Validation Accuracy(%)")
plt.show()
```



Trying out the best model on MNIST Dataset

Loading Dataset

```
In [21]: (train, train_labels), (test, test_labels) = tf.keras.datasets.mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```

```
In [22]: def draw_matrix(model, train, test, train_labels, test_labels):
    predictions_test = []
    for item in test:
        predictions_test.append(np.argmax(model.predict(item.reshape(-1,28,28,1)))))

    cm_test = sklearn.metrics.confusion_matrix(np.argmax(test_labels, axis = 1), predictions_test)
    sns.heatmap(cm_test, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues_r');
    plt.ylabel('Actual label');
    plt.xlabel('Predicted label');
    all_sample_title = 'Confusion Matrix (Test samples)'
    plt.title(all_sample_title)
    plt.show()
    print(classification_report(np.argmax(test_labels, axis = 1), predictions_test))

def draw(History, model, train, test, train_labels, test_labels):
    draw_graph(History)
    draw_matrix(model, train, test, train_labels, test_labels)
```

Preparing the images for feeding into CNN

```
In [23]: train = train.reshape((-1,28,28,1))
test = test.reshape((-1,28,28,1))
```

Categorical (One-hot) Encoding

```
In [24]: from keras.utils import np_utils
train_labels = np_utils.to_categorical(train_labels, 10)
test_labels = np_utils.to_categorical(test_labels, 10)
```

Model

```
In [25]: leakyrelu = tf.keras.layers.LeakyReLU(alpha=0.01)
model18 = Sequential()

model18.add(Conv2D(filters = 16, kernel_size = (3,3), padding = 'Same', activation = leakyrelu, input_shape = (28,28,1)))
```

```
model18.add(MaxPooling2D(pool_size=(2, 2)))
model18.add(Dropout(0.1))

model18.add(Conv2D(filters = 32, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model18.add(MaxPooling2D(pool_size=(2, 2)))
model18.add(Dropout(0.1))

model18.add(Conv2D(filters = 64, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model18.add(MaxPooling2D(pool_size=(2, 2)))
model18.add(Dropout(0.1))

model18.add(Conv2D(filters = 128, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model18.add(MaxPooling2D(pool_size=(2, 2)))
model18.add(Dropout(0.1))

model18.add(Conv2D(filters = 256, kernel_size = (5,5),padding = 'Same',activation = leakyrelu))
model18.add(Dropout(0.1))

model18.add(Flatten())
model18.add(Dense(128, activation = leakyrelu))
model18.add(Dense(256, activation = leakyrelu))
model18.add(Dense(10, activation = 'softmax'))

model18.compile(optimizer=Adam(),loss='categorical_crossentropy',metrics=[ 'accuracy'])
model18.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0
dropout (Dropout)	(None, 14, 14, 16)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	12832
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0
dropout_1 (Dropout)	(None, 7, 7, 32)	0
conv2d_2 (Conv2D)	(None, 7, 7, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 64)	0
dropout_2 (Dropout)	(None, 3, 3, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 128)	204928
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 128)	0
dropout_3 (Dropout)	(None, 1, 1, 128)	0
conv2d_4 (Conv2D)	(None, 1, 1, 256)	819456
dropout_4 (Dropout)	(None, 1, 1, 256)	0
conv2d_5 (Conv2D)	(None, 1, 1, 512)	3277312
dropout_5 (Dropout)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 128)	65664

dense_1 (Dense)	(None, 256)	33024
dense_2 (Dense)	(None, 10)	2570

```
=====
Total params: 4,467,210
Trainable params: 4,467,210
Non-trainable params: 0
```

```
In [26]: History18 = model18.fit(train,train_labels, batch_size=128, epochs = 10, validation_split = 0.1)
```

```
Epoch 1/10
422/422 [=====] - 16s 15ms/step - loss: 0.3383 - accuracy: 0.8872 - val_loss: 0.0690 - val_accuracy: 0.9810
Epoch 2/10
422/422 [=====] - 6s 15ms/step - loss: 0.0909 - accuracy: 0.9754 - val_loss: 0.0508 - val_accuracy: 0.9862
Epoch 3/10
422/422 [=====] - 6s 14ms/step - loss: 0.0670 - accuracy: 0.9819 - val_loss: 0.0464 - val_accuracy: 0.9878
Epoch 4/10
422/422 [=====] - 7s 16ms/step - loss: 0.0563 - accuracy: 0.9854 - val_loss: 0.0369 - val_accuracy: 0.9913
Epoch 5/10
422/422 [=====] - 6s 14ms/step - loss: 0.0514 - accuracy: 0.9865 - val_loss: 0.0449 - val_accuracy: 0.9890
Epoch 6/10
422/422 [=====] - 6s 14ms/step - loss: 0.0452 - accuracy: 0.9883 - val_loss: 0.0499 - val_accuracy: 0.9897
Epoch 7/10
422/422 [=====] - 6s 14ms/step - loss: 0.0405 - accuracy: 0.9896 - val_loss: 0.0392 - val_accuracy: 0.9923
Epoch 8/10
422/422 [=====] - 6s 15ms/step - loss: 0.0396 - accuracy: 0.9895 - val_loss: 0.0451 - val_accuracy: 0.9902
Epoch 9/10
422/422 [=====] - 6s 14ms/step - loss: 0.0348 - accuracy: 0.9906 - val_loss: 0.0499 - val_accuracy: 0.9902
Epoch 10/10
422/422 [=====] - 6s 14ms/step - loss: 0.0330 - accuracy: 0.9911 - val_loss: 0.0378 - val_accuracy: 0.9912
```

Making predictions on 5 test images

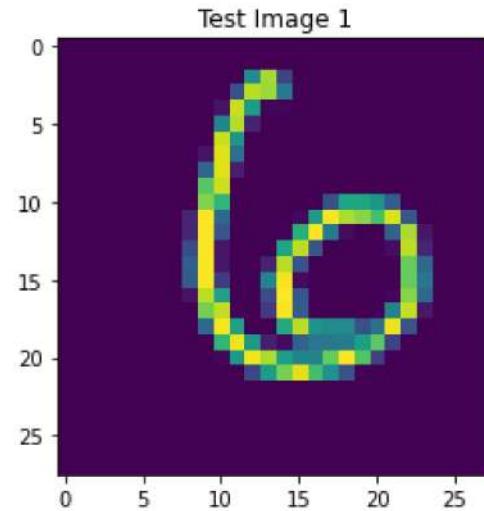
```
In [27]: results = model18.predict(test)

def test_inference(id, results):
    img = test[id]
    img = img.reshape((28,28))
    plt.imshow(img)
    plt.title("Test Image 1")
    print("True class:", np.argmax(test_labels[id]))
    print("Predicted class:", np.argmax(results[id]))
    plt.show()
```

313/313 [=====] - 2s 4ms/step

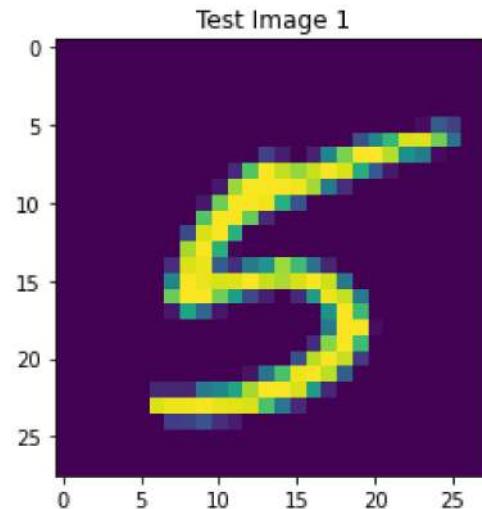
```
In [28]: test_inference(100, results)
```

True class: 6
Predicted class: 6



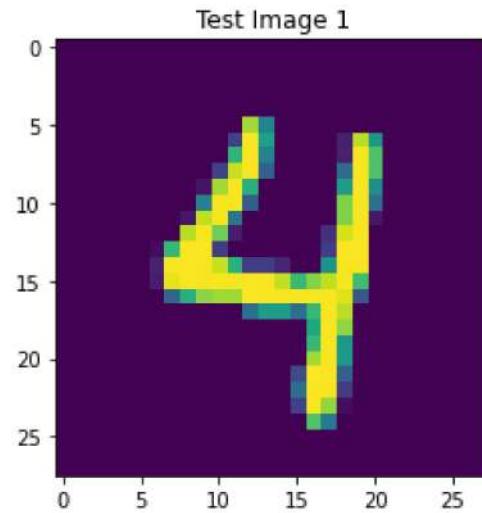
```
In [29]: test_inference(52, results)
```

True class: 5
Predicted class: 5



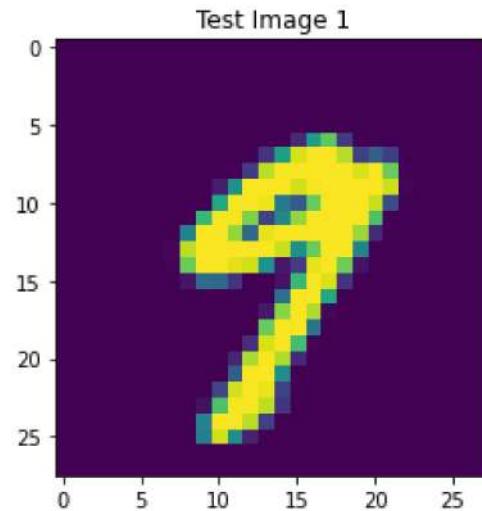
```
In [30]: test_inference(27, results)
```

True class: 4
Predicted class: 4



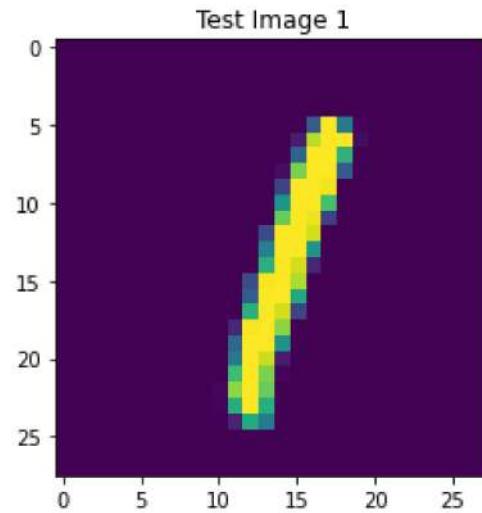
```
In [31]: test_inference(20, results)
```

True class: 9
Predicted class: 9



```
In [32]: test_inference(5, results)
```

```
True class: 1  
Predicted class: 1
```



Testing on 5 handwritten images

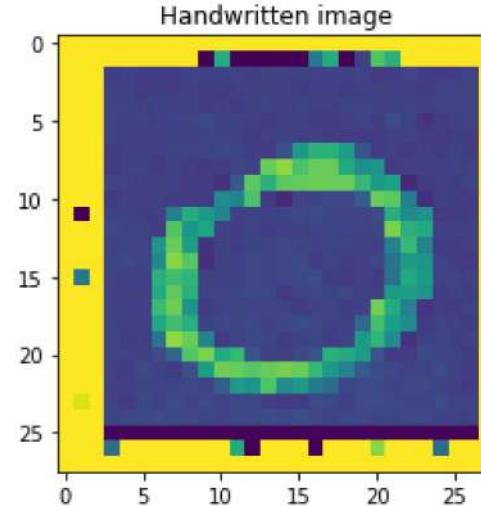
```
In [100...]  
def load_image(path):  
    img = cv2.imread(path)  
    return img
```

```
def inference(img):
    img_size = (28,28)
    img = np.array(img)
    img = cv2.resize(img, img_size)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    test_img = img.reshape((-1,28,28,1))
    img_class = np.argmax(model18.predict(test_img), axis=-1)
    prediction = img_class[0]
    classname = img_class[0]
    print("Predicted Class: " + (str)(classname))
    img = img.reshape((28,28))
    plt.imshow(img)
    plt.title("Handwritten image")
    plt.show()
```

In [101]: BASE_PATH = "/content/drive/My Drive/ML_DRIVE/Assign_6/"

In [104]: path = BASE_PATH + "0.png"
img1 = load_image(path)
inference(img1)

1/1 [=====] - 0s 15ms/step
Predicted Class: 0



In []: path = BASE_PATH + "1.png"
img1 = load_image(path)

```
inference(img1)
```

```
In [ ]: path = BASE_PATH + "2.png"
img1 = load_image(path)
inference(img1)
```

```
In [ ]: path = BASE_PATH + "3.png"
img1 = load_image(path)
inference(img1)
```

```
In [ ]: path = BASE_PATH + "6.png"
img1 = load_image(path)
inference(img1)
```

```
In [ ]:
```